

Training Industry-scale GNNs with GiGL

KDD 2025 Hands-On tutorial

Presenters: Yozén Liu, Tong Zhao, Matthew Kolodner, Kyle Montemayor, Shubham Vij, Neil Shah



https://github.com/Snapchat/GiGL/tree/main/examples/tutorial/KDD_2025

Hands-on Tutorial website



Overview

- /01 **Intro**
- /02 **GNNs and their scale challenges**
- /03 **Overview of GiGL**
- /04 **Hands-on with GiGL - Training and inference with industry-scale GNNs**
- Break (30m)**
- /05 **Hands-on with GiGL - Customization**

SNAPCHAT 



Intro



Who are we?

User Modeling and Personalization: A group of scientists and engineers at Snapchat advancing personalization-related algorithms and applications.

We do both fundamental research and support internal technical transfer.

Focus areas:

- graph learning
- classical and generative recommendation systems
- efficient ML at scale
- new applications

Roles at Snap



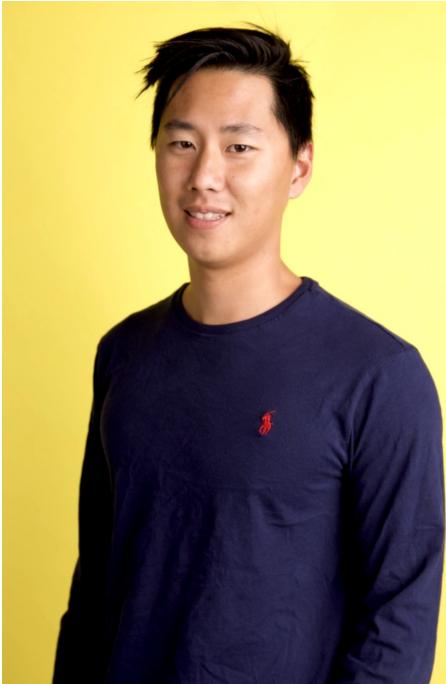
Internships



Who are we?



Shubham



Yozan



Matt



Kyle



Neil





Why are we here?

We will showcase **GiGL (Gigantic Graph Learning)**, a library we built at Snapchat to make scaling up and iterating on graph neural networks easy.

By the end of the tutorial, you will...

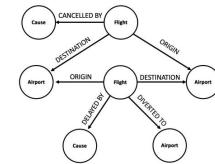
- Understand fundamentals of GNNs, scaling challenges, and GiGL as a solution
- Gain hands-on experience using GiGL components
- Be able to comfortably train your own large-scale GNNs with GiGL

GNNs and their Scale Challenges



Graphs and why they matter

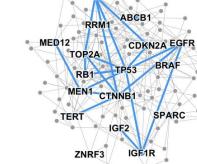
Graphs are a universal language to describe relationships between entities.



Event Graphs



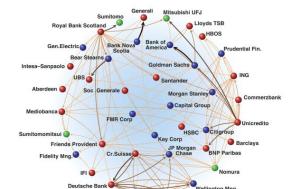
Computer Networks



Disease Pathways



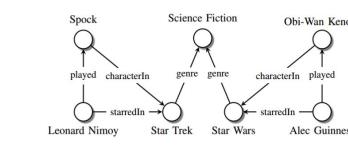
Social Networks



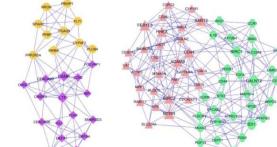
Economic Networks



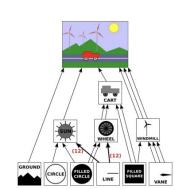
Communication Networks



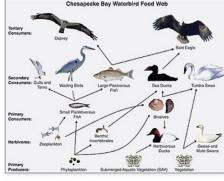
Knowledge Graphs



Regulatory Networks



Scene Graphs



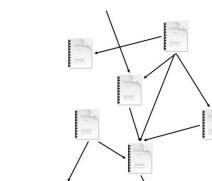
Food Webs



Particle Networks



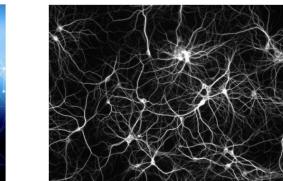
Underground Networks



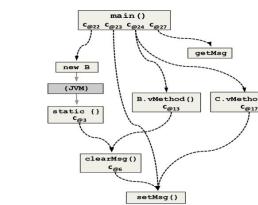
Citation Networks



Internet

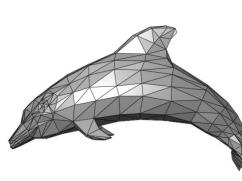


Networks of Neurons



Code Graphs

Molecules

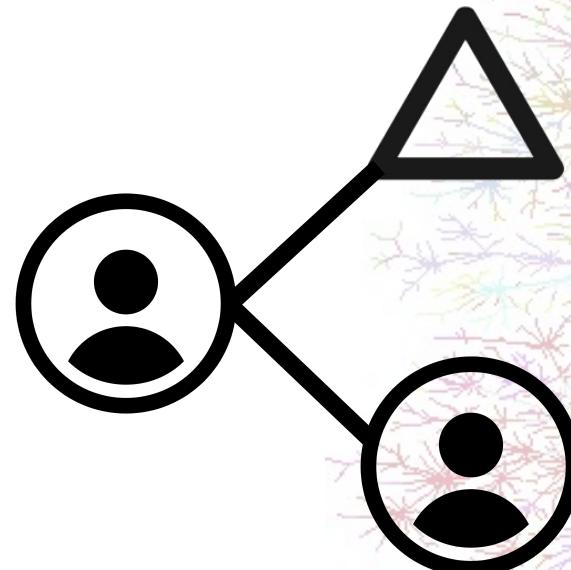


3D Shapes

This is immensely valuable context in making predictions.



What is **gigantic**, really?



Friendships



Snapping



Chatting



Story viewing



Publisher
content



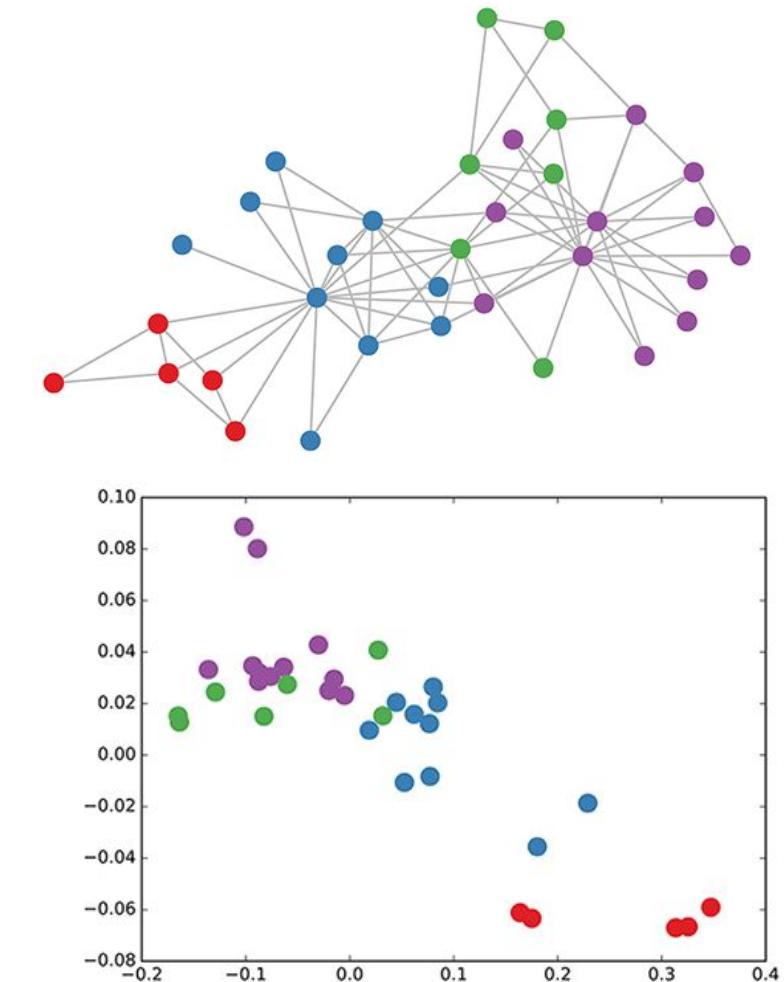
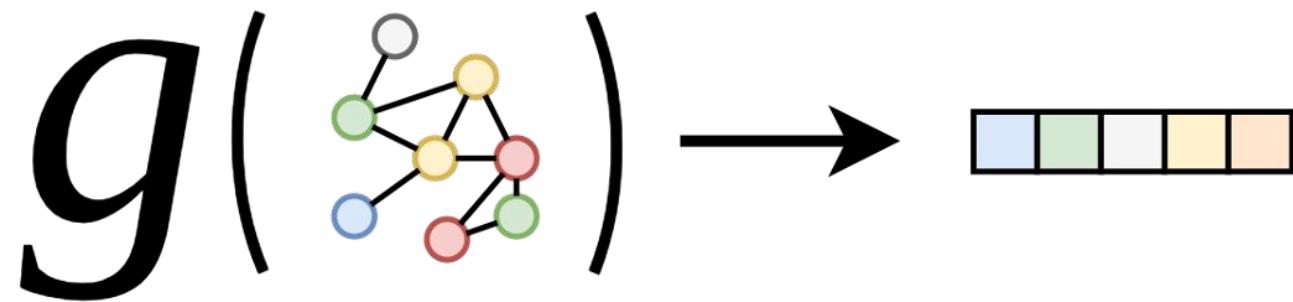
Lenses

0(billion) nodes, 0(100B) edges



What are Graph Neural Networks (GNNs)?

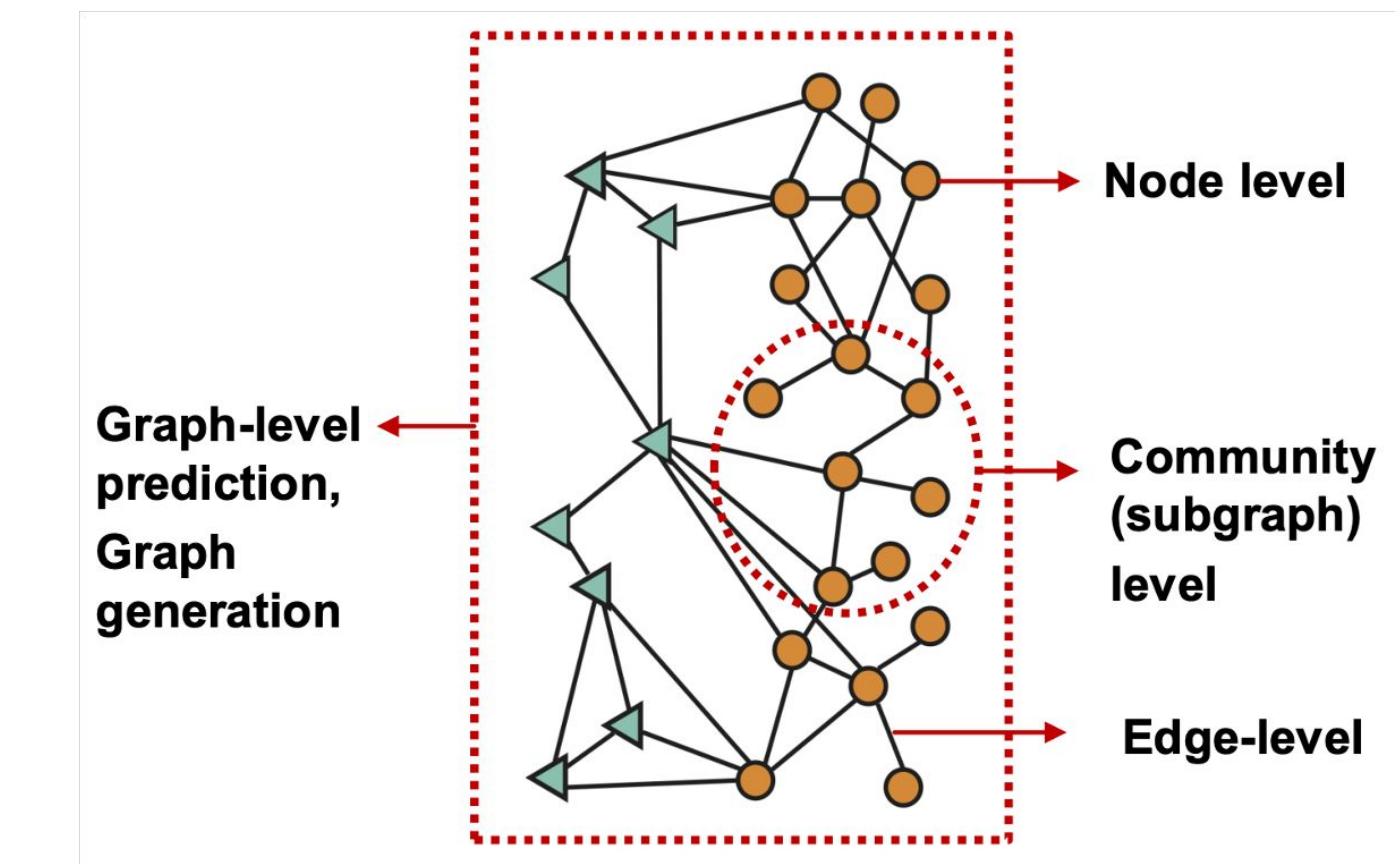
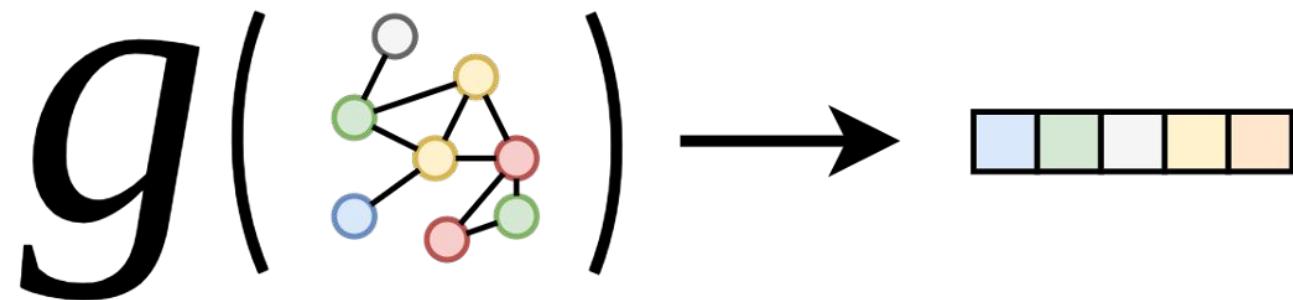
GNNs are a modern neural architecture designed to model graph data





What can we do with them?

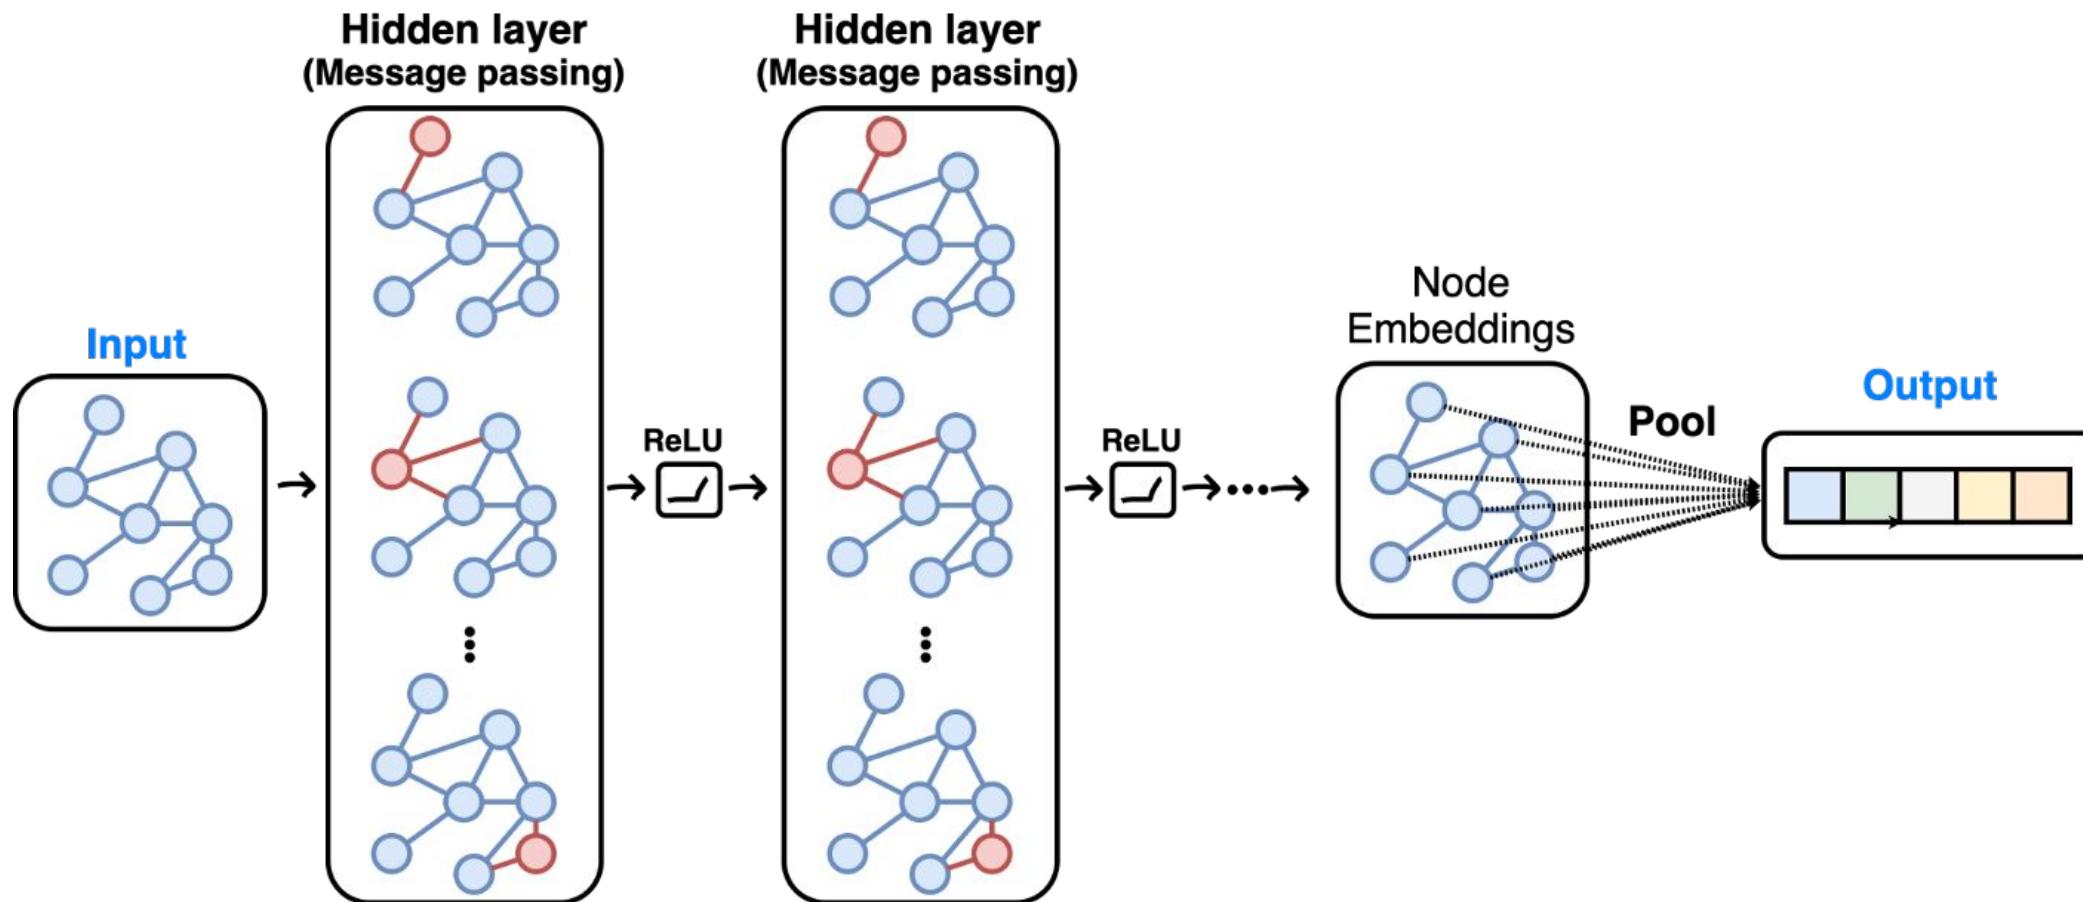
GNNs are a modern neural architecture designed to model graph data





How do they work?

Key Idea: stack “message passing” layers with nonlinearities





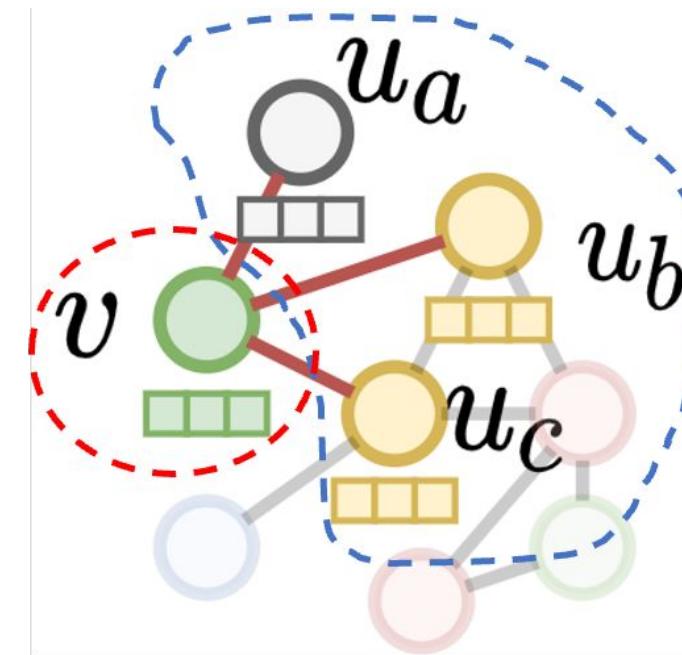
How do they work, **really**?

$$h_v^{(t)} = \text{AGG}^{(t)} \left(h_v^{(t-1)}, \left\{ \text{MSG}^{(t)}(h_u^{(t-1)}) \mid u \in \mathcal{N}(v) \right\} \right)$$

█ Node emb. █ Neighbor messages

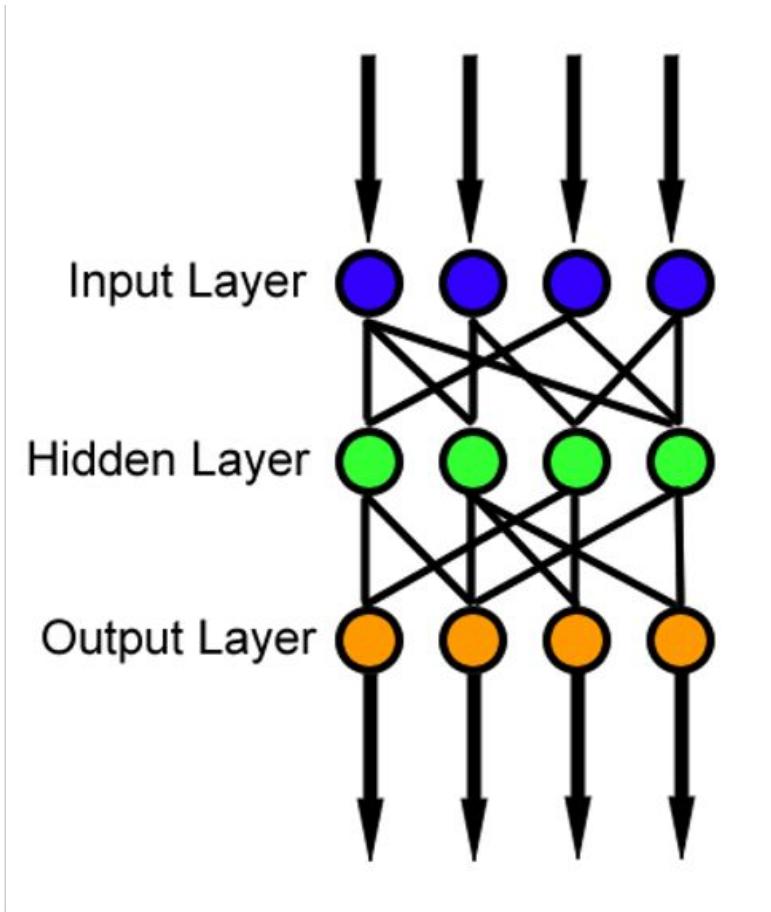
$$H^{(t)} = \sigma \left(A H^{(t-1)} W^{(t-1)} \right)$$

where $H^{(0)} = X$

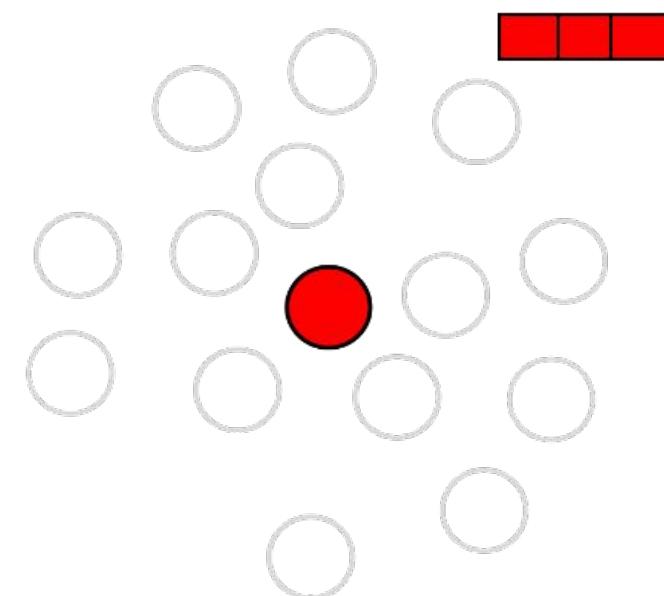




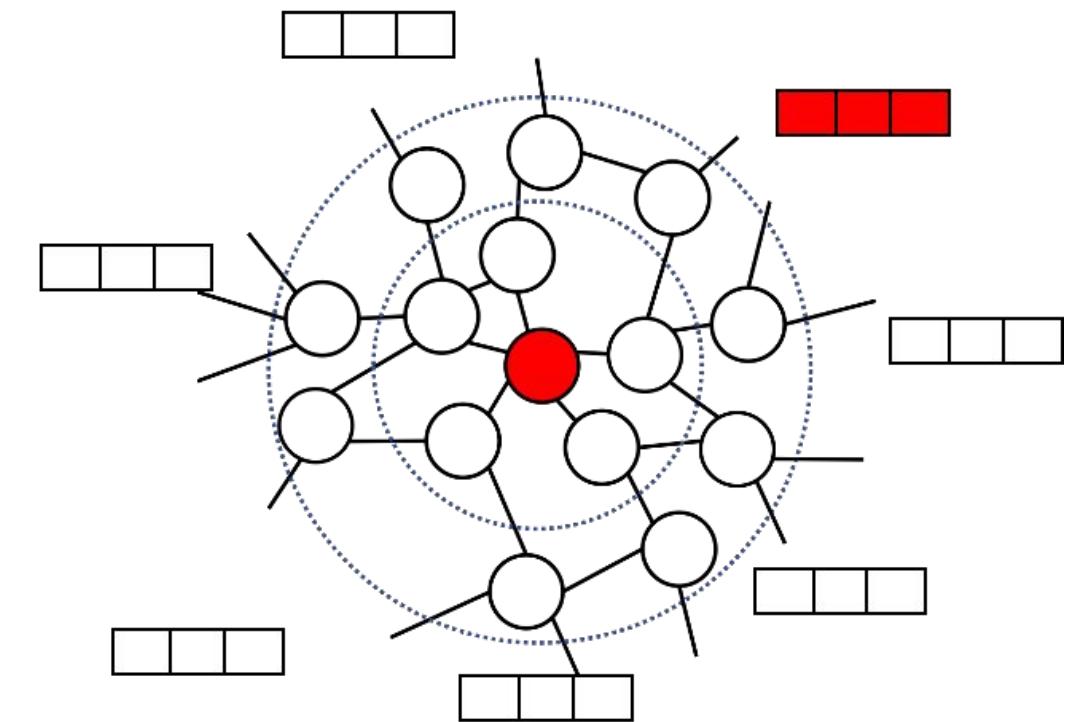
Connections to Multi-Layer Perceptrons



Multi-layer perceptrons
(tabular machine learning)



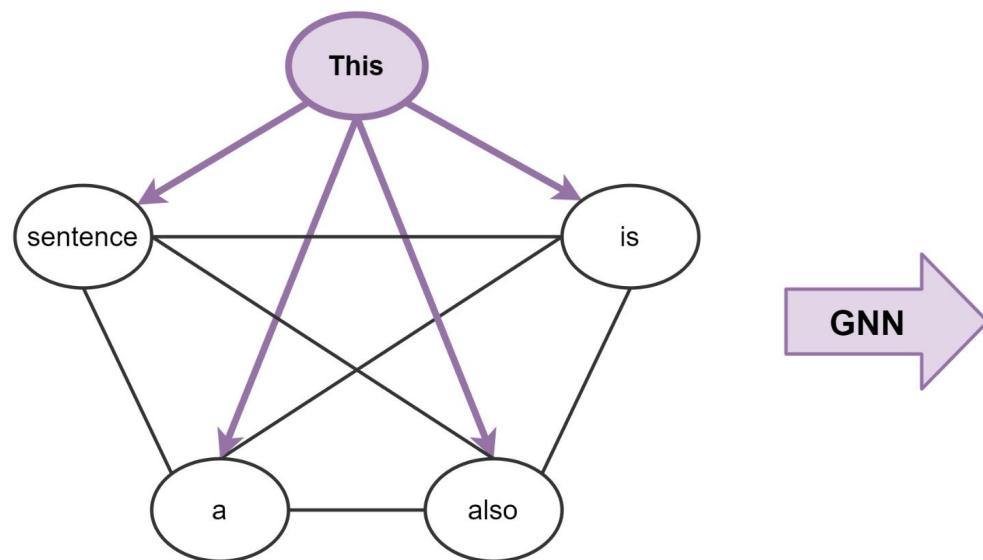
Tabular setting



Relational setting



Connections to Transformers



Translation?
Sentiment?
Next word?
Part-of-speech tags?

Transformers are Graph Neural Networks

Chaitanya Joshi

Feb 12, 2020

[Project](#) [Project](#) [Slides](#) [Medium](#)

Engineer friends often ask me: Graph Deep Learning sounds great, but are there any big commercial success stories? Is it being deployed in practical applications?

Besides the obvious ones—recommendation systems at [Pinterest](#), [Alibaba](#) and [Twitter](#)—a slightly nuanced success story is the [Transformer architecture](#), which has [taken the NLP industry by storm](#).

Through this post, I want to establish links between [Graph Neural Networks \(GNNs\)](#) and Transformers. I'll talk about the intuitions behind model architectures in the NLP and GNN communities, make connections using equations and figures, and discuss how we could work together to drive progress.

Let's start by talking about the purpose of model architectures—*representation learning*.



GNNs are widely successful in practice

Uber AI, Engineering

Food Discovery with Uber Eats: Using Graph Learning to Power Recommendations

LinkSAGE: Optimizing Job Matching Using Graph Neural Networks

Ping Liu, Haichao Wei, Xiaochen Hou, Jianqiang Shen, Shihai He, Kay Qianqi Shen, Zhujun Chen, Fedor Borisyuk, Daniel Hewlett, Liang Wu, Srikant Veeraraghavan, Alex Tsun, Chengming Jiang, Wenjing Zhang
LinkedIn Corporation
Mountain View, CA, USA
 {piliu,hawei,xiahou,jershenshe1,qishen,zhuchen,fborisyuk,dhewlett,liawu,atsun,cjiang,wzhang}@linkedin.com

Personalized Audiobook Recommendations at Spotify Through Graph Neural Networks

Marco De Nadai^{1*}, Francesco Fabbri^{1*}, Paul Gigioli¹, Alice Wang¹, Ang Li¹, Fabrizio Silvestri^{1,2}, Laura Kim¹, Shawn Lin¹, Vladan Radosavljevic¹, Sandeep Ghosh¹, David Nyhan¹, Hugues Bouchard¹, Mounia Lalmas-Roelleke¹, Andreas Damianou¹

¹Spotify, Denmark, Spain, UK, USA

²Sapienza University of Rome, Italy

PinSage: A new graph convolutional neural network for web-scale recommender systems

ETA Prediction with Graph Neural Networks in Google Maps

Austin Derrow-Pinion¹, Jennifer She¹, David Wong^{2*}, Oliver Lange³, Todd Hester^{4*}, Luis Perez^{5*}, Marc Nunkesser³, Seongjae Lee³, Xueying Guo³, Brett Wiltshire¹, Peter W. Battaglia¹, Vishal Gupta¹, Ang Li¹, Zhongwen Xu^{6*}, Alvaro Sanchez-Gonzalez¹, Yujia Li¹, Petar Veličković¹
¹DeepMind ²Waymo ³Google ⁴Amazon ⁵Facebook AI ⁶Sea AI Lab *work done while at DeepMind
 {derrowap,jenshe,wongda,petary}@google.com

Embedding Based Retrieval in Friend Recommendation

Jiahui Shi, Vivek Chaurasiya, Yozen Liu, Shubham Vij, Yan Wu, Satya Kanduri
 Neil Shah, Peicheng Yu, Nik Srivastava, Lei Shi, Ganesh Venkataraman, Jun Yu
 {jshi3,vchaurasiya,yliu2,svij,ywu,satya.kanduri,nshah,peicheng.yu,nsrivastava,lshi3,gvenkataraman,jyu3}@snap.com
 Snap Inc.
 Santa Monica, CA USA

...especially at Snap!

GiGL: Large-Scale Graph Neural Networks at Snapchat

Tong Zhao*, Yozen Liu*, Matthew Kolodner, Kyle Montemayor, Elham Ghazizadeh[†], Ankit Batra[†]
Zihao Fan, Xiaobin Gao, Xuan Guo, Jiwen Ren, Serim Park, Peicheng Yu, Jun Yu
Shubham Vij*, Neil Shah*

Snap Inc., USA

*{tong,yliu2,svij,nshah}@snap.com

ABSTRACT

Recent advances in graph machine learning (ML) with the introduction of Graph Neural Networks (GNNs) have led to a widespread interest in applying these approaches to business applications at scale. GNNs enable differentiable end-to-end (E2E) learning of model parameters given graph structure which enables optimization towards popular node, edge (link) and graph-level tasks. While the research innovation in new GNN layers and training strategies has been rapid, industrial adoption and utility of GNNs has lagged considerably due to the unique scale challenges that large-scale graph ML problems create. In this work, we share our approach to training, inference, and utilization of GNNs at Snapchat. To this end, we present GiGL (Gigantic Graph Learning), an open-source library to enable large-scale distributed graph ML to the benefit of researchers, ML engineers, and practitioners. We use GiGL internally at Snapchat to manage the heavy lifting of GNN workflows, including graph data preprocessing from relational DBs, subgraph sampling, distributed

1 INTRODUCTION

Graphs are ubiquitous, representing a wide variety of real-world data across domains such as social networks and recommendation systems. In recent years, Graph Neural Networks (GNNs) have emerged as powerful tools for learning from graph data [30, 51, 63].

Nevertheless, when deploying GNNs in industrial contexts, such as large-scale user modeling and social recommendation systems, scalability poses a significant challenge. Most commonly used libraries in the research community assume that graph topology, node and edge features, can fit easily within the CPU (or even GPU) memory of a single machine. However, just storing a graph with 100 billion edges as an example (not uncommon in practice [6]) requires 800 GB of memory assuming 32-bit integer IDs, even before accounting for node and edge features or the need to represent graphs with more nodes than 32-bit limits. Such limitations make it difficult to scale GNN models in real-world, billion-scale graphs, which we routinely handle at Snapchat.

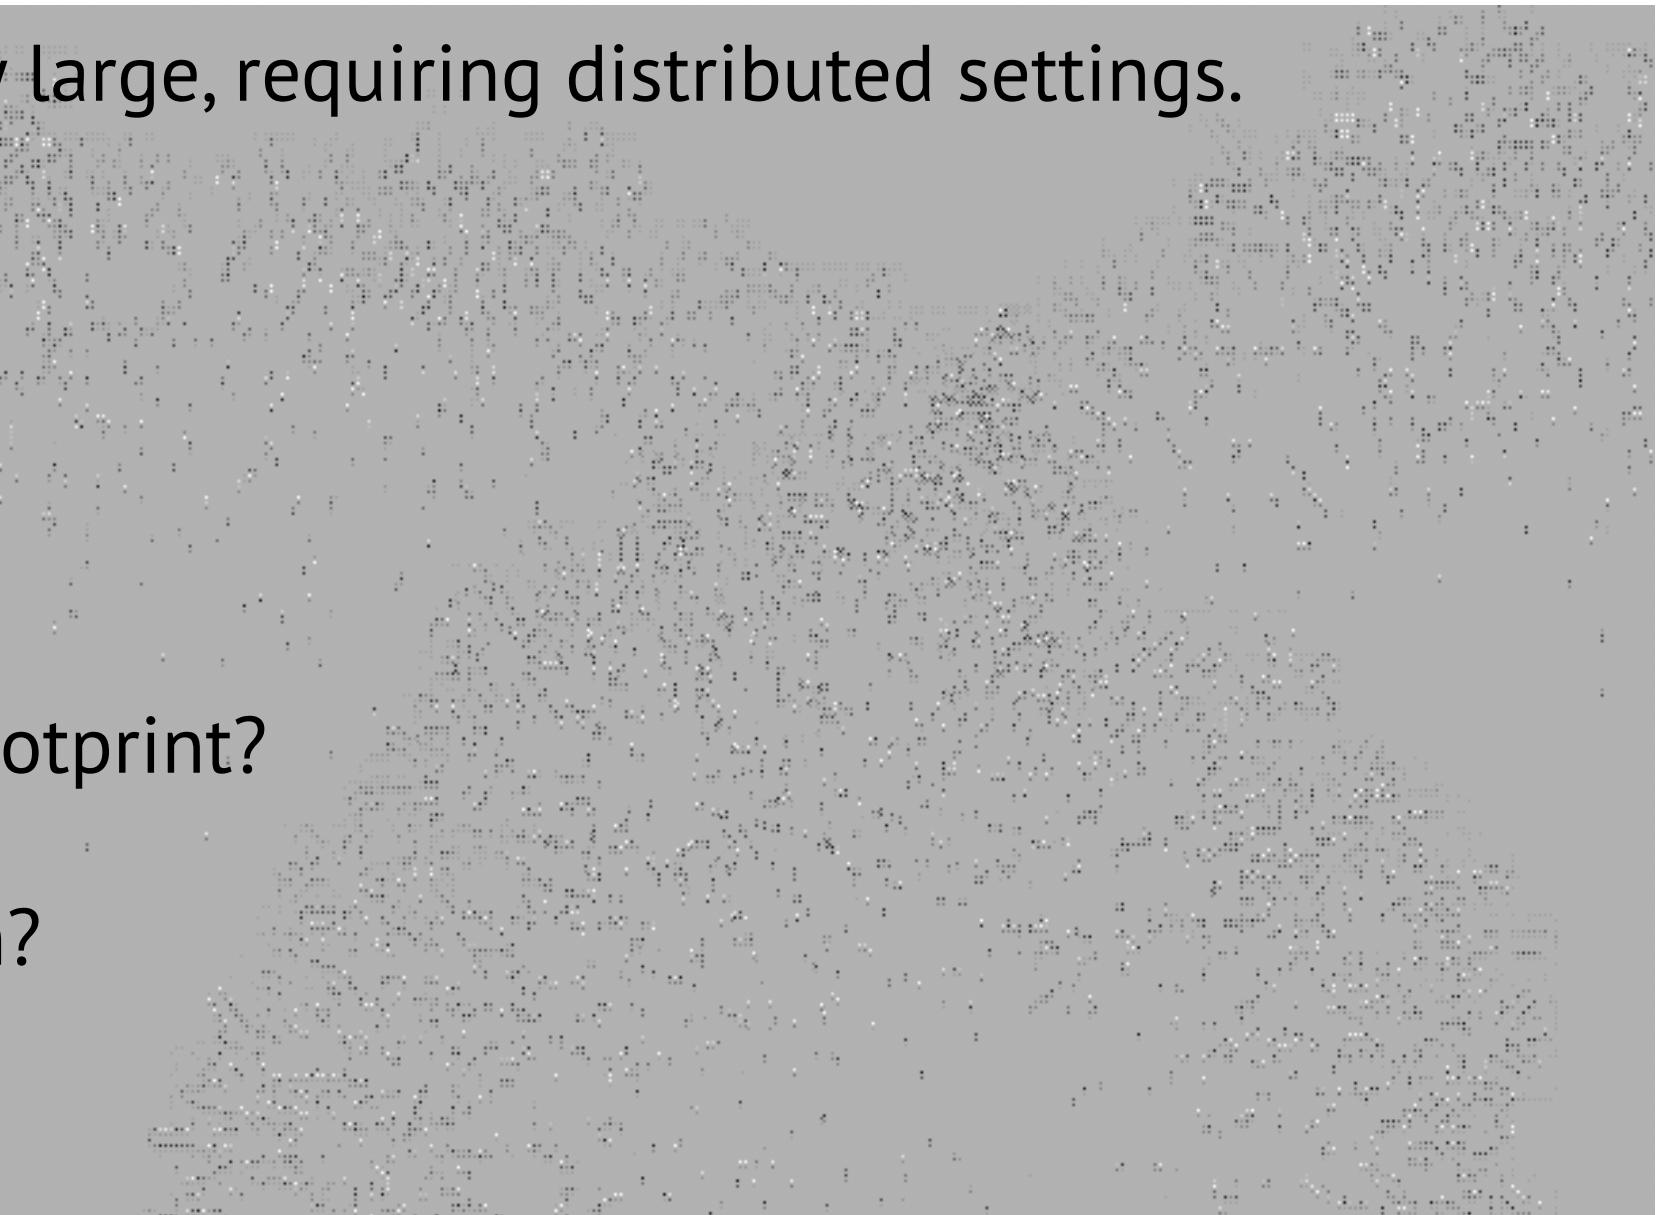




Practical Challenges in Scaling GNNs

Graphs in practice can be extremely large, requiring distributed settings.

- How to represent graph data?
- How to manage memory usage?
- How to minimize computation footprint?
- How to minimize communication?

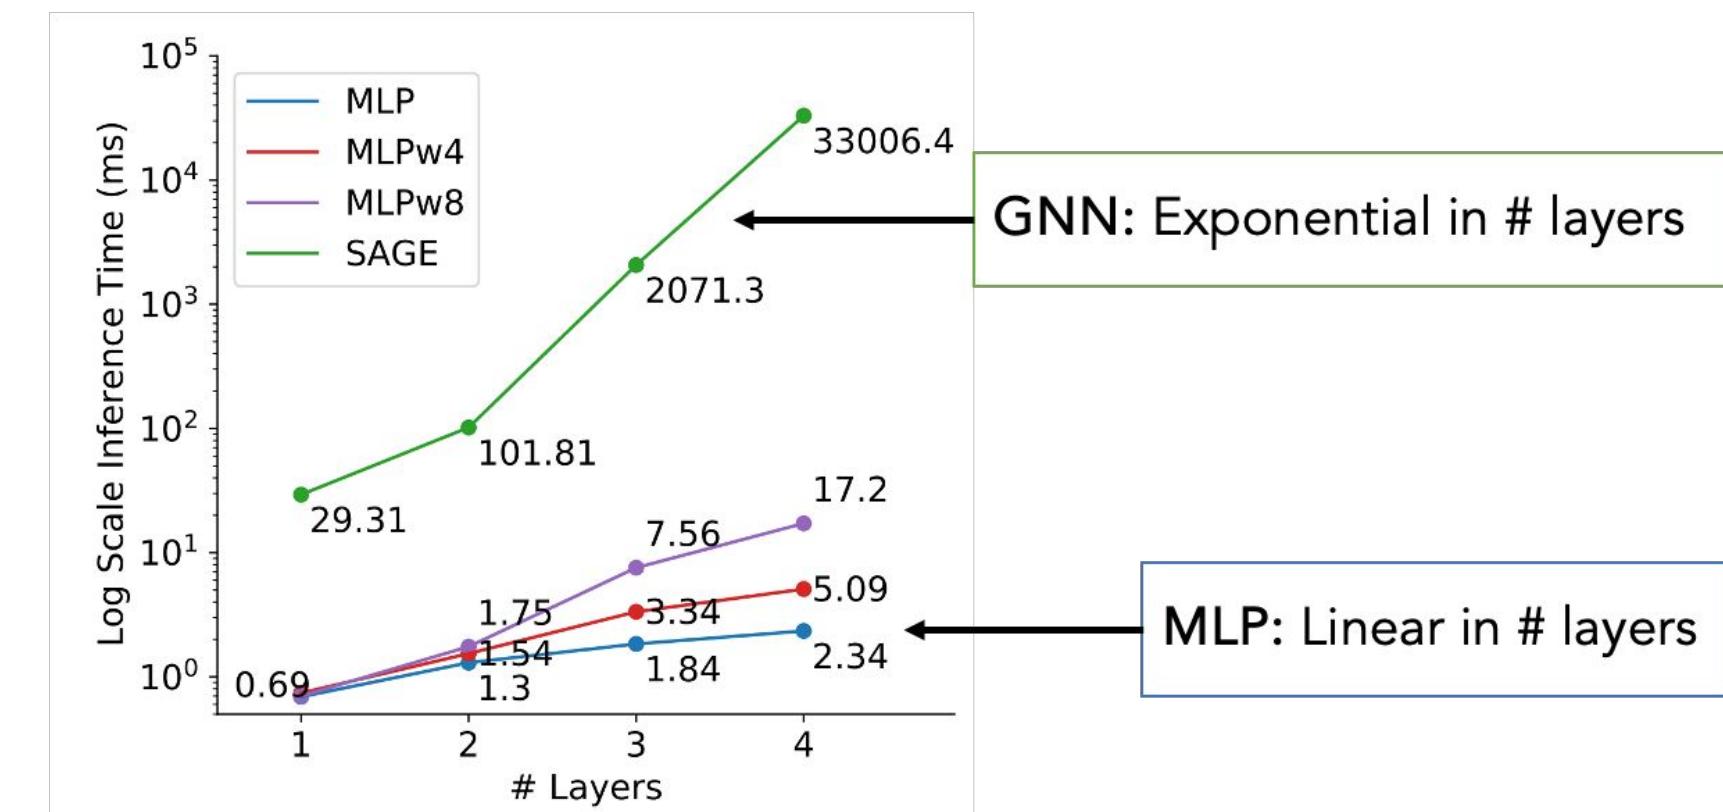


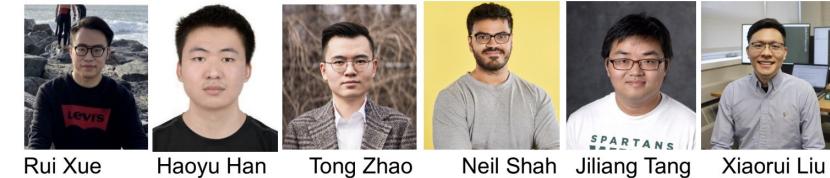


Practical Challenges in Scaling GNNs

Neighbor explosion in GNNs creates scaling overheads.

Operation	OGB-arXiv		
	Forward	Backward	Total
$Z^l = W^l H^{l-1}$	0.32	1.09	1.42
$H^l = AZ^l$	1.09	1028.08	1029.17 724×





8:30 am – 12:30 pm, February 21th, AAAI 2024



GNN Scaling: Research

Approximations: node and layer-wise sampling, historical embeddings

Decoupling: pre-computing and post-computing with the graph

Distributed processing: efficient algorithms for parallelization and tabularization

Training paradigm: lazy propagation, alternating opt, pre-training, coarsening

Inference paradigm: cross-model distillation

Data-centric: graph condensation, subgraph sketching

GNN Scaling: Practice

In our experience, we want tooling which enables...

- Flexible model specification and iteration
- Familiar imperatives for ML practitioners (e.g. PyG)
- Orchestration for batch training and inference pipelines
- Elastic horizontal scaling to large graphs
- Minimal cost and e2e latency





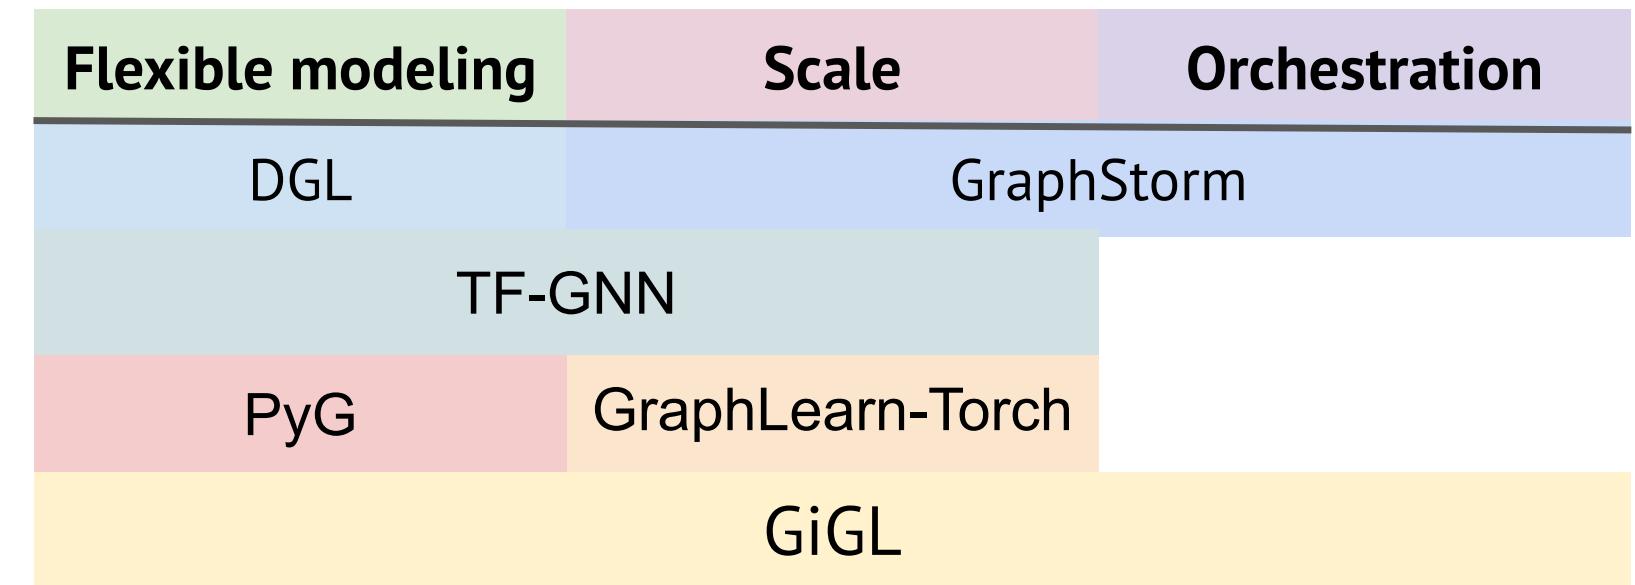
Overview of GiGL

What is GiGL?

- GiGL is a graph learning library built for training **industry-scale graphs**
- GiGL provides **end-to-end pipeline and orchestration** handling all pre/post-processing, training and inference steps
- GiGL provides easy onboarding, extensive code customization, and is compatible with popular graph learning frameworks like **PyTorch Geometric**.

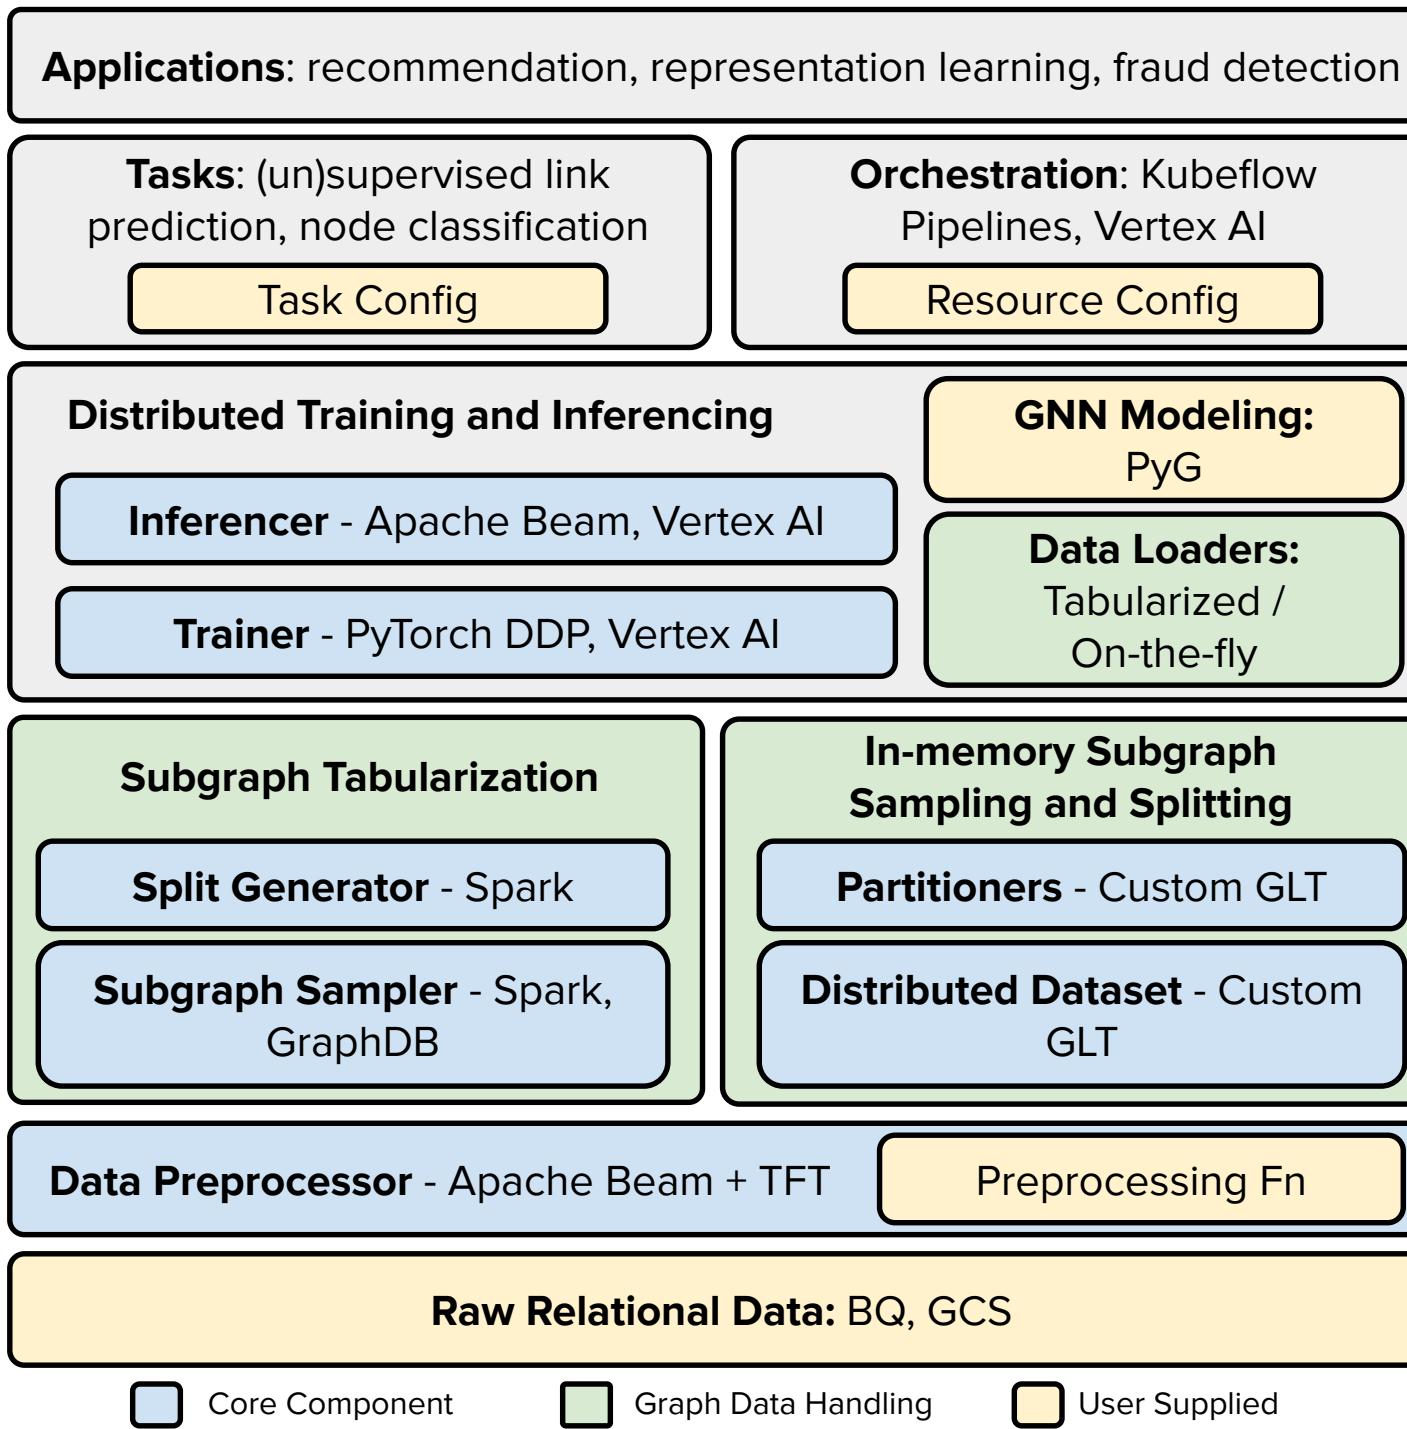


Why build GiGL?



- Existing Graph ML tooling does not support our requirements on -
 - The scale of graph we operate on
 - End-to-end orchestration that seamlessly integrates into our ecosystem
 - Flexible modeling and subgraph sampling customization

Bird's Eye View



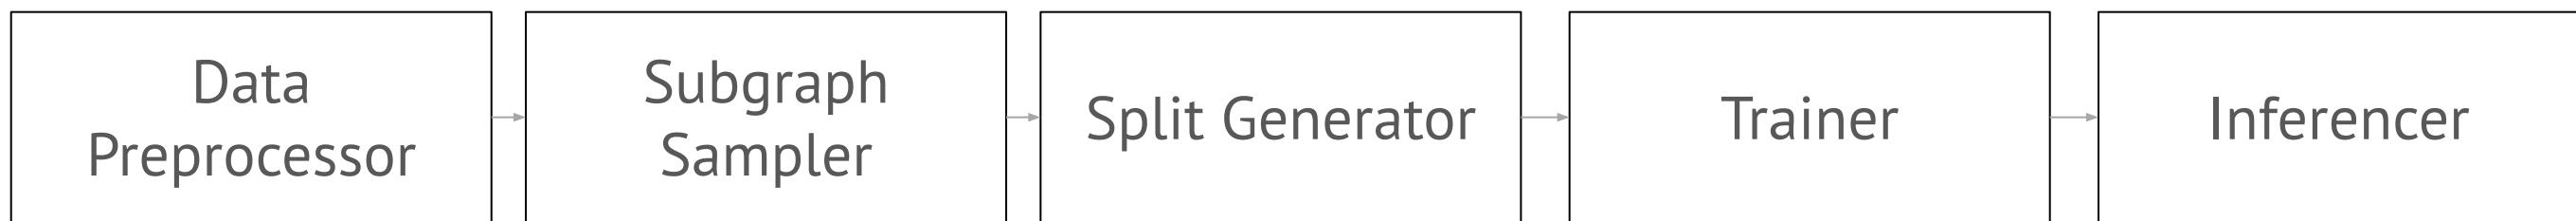
Overview of GiGL

Solving scaling challenges

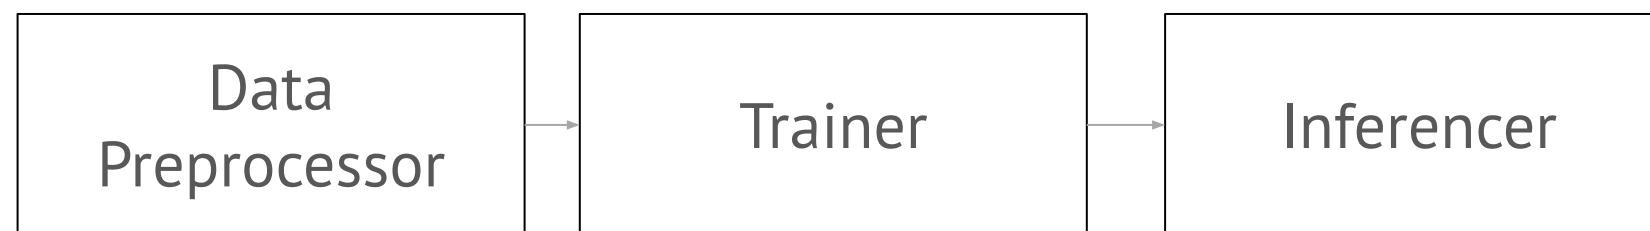


How does GiGL solve scale? - Two paradigms

Subgraph Tabularization (Powered by Spark)

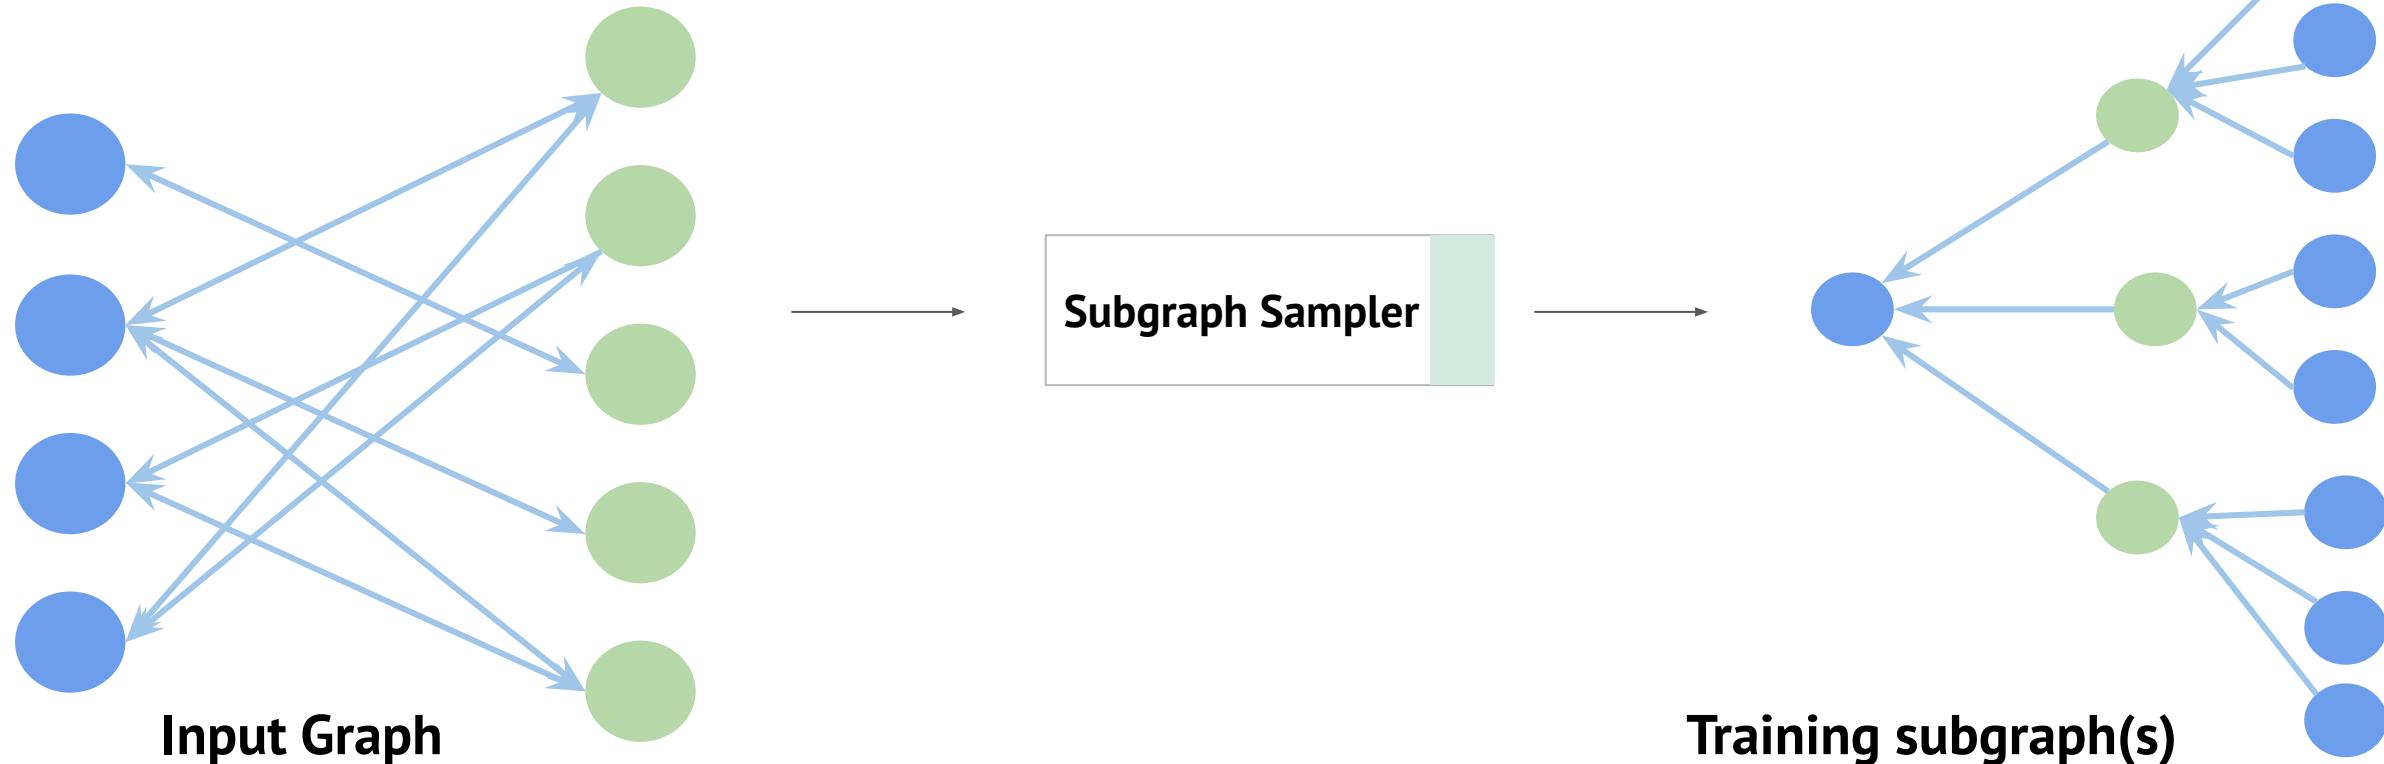


In-memory Subgraph Sampling (Powered by GraphLearn-for-PyTorch [GLT])





Subgraph Tabularization



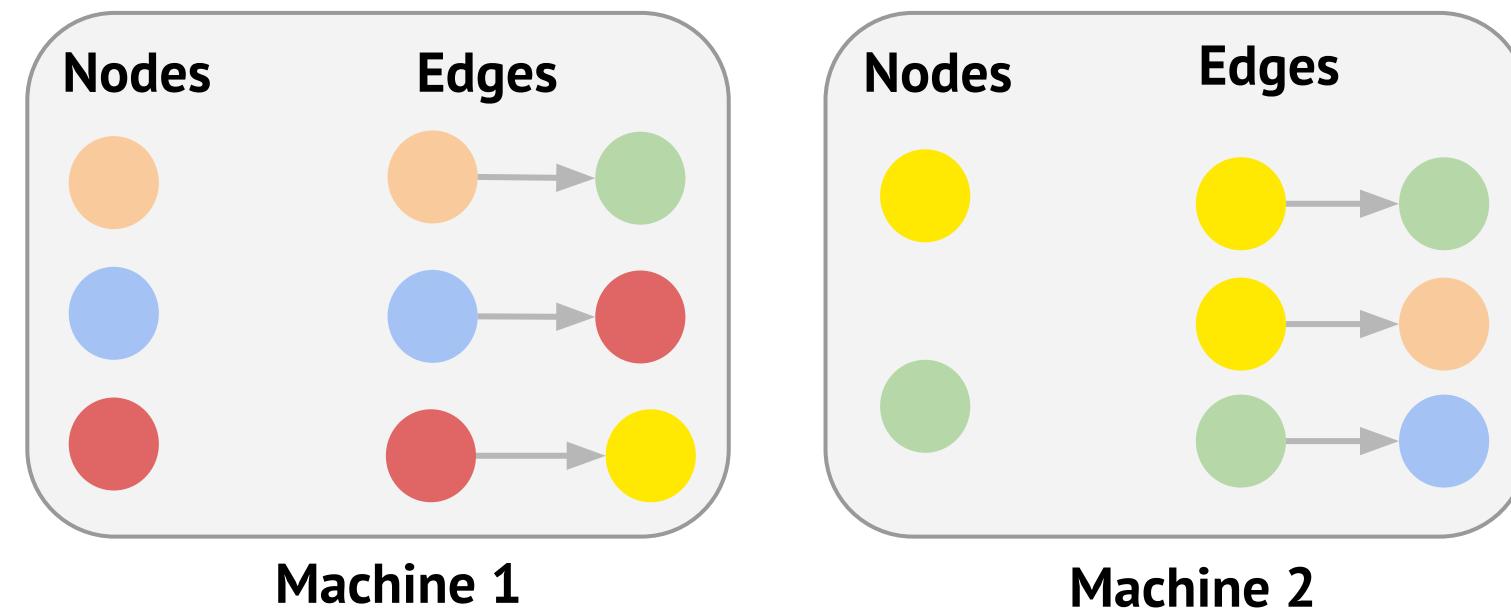
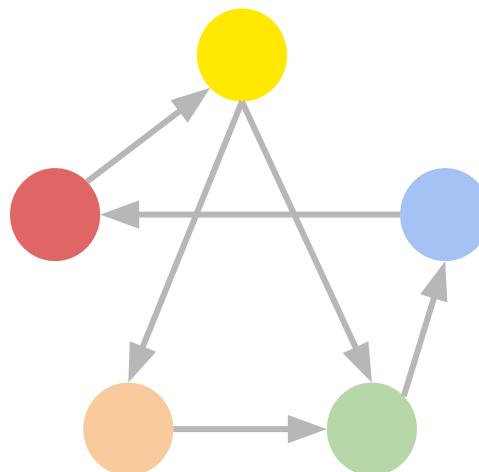
Enables training GNNs at scale by generating training subgraphs stored on disk in a distributed processing (Spark) job prior to training



In-memory Subgraph Sampling

Graph is stored in-memory

- Nodes & edges partitioned into multiple machines
- Training subgraphs sampled during training process



Tabularized vs In-memory - When to use which?

Tabularization

- Offline batch jobs
- Multiple tuning/experimental runs benefit from pre-sampled training subgraphs

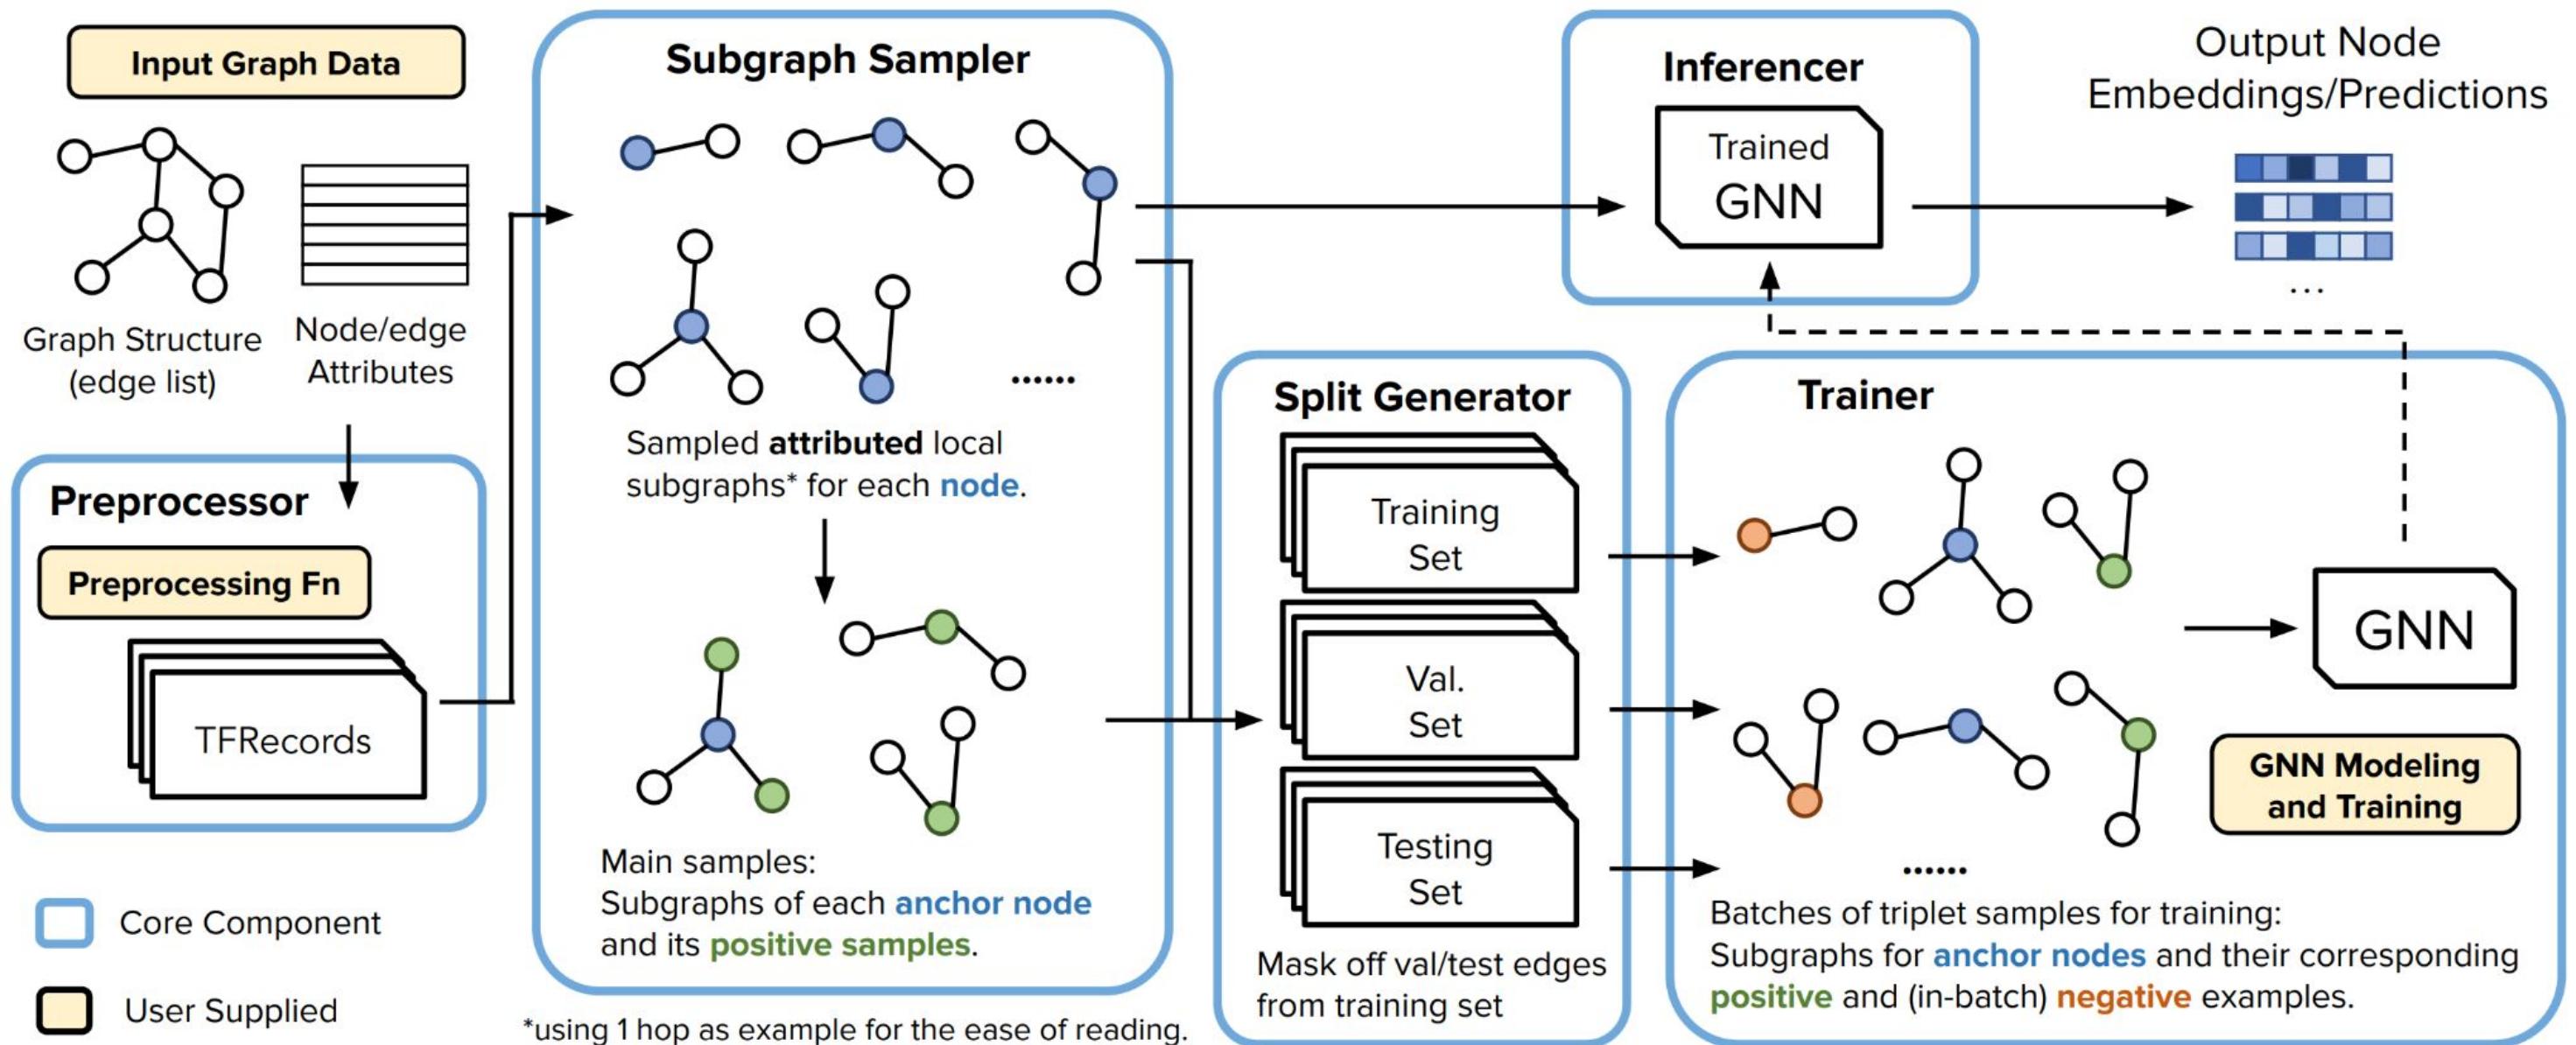
In-memory

- Offline and Online workloads and real-time inference
- Needs machines with large memory
- Efficient in-memory subgraph sampling

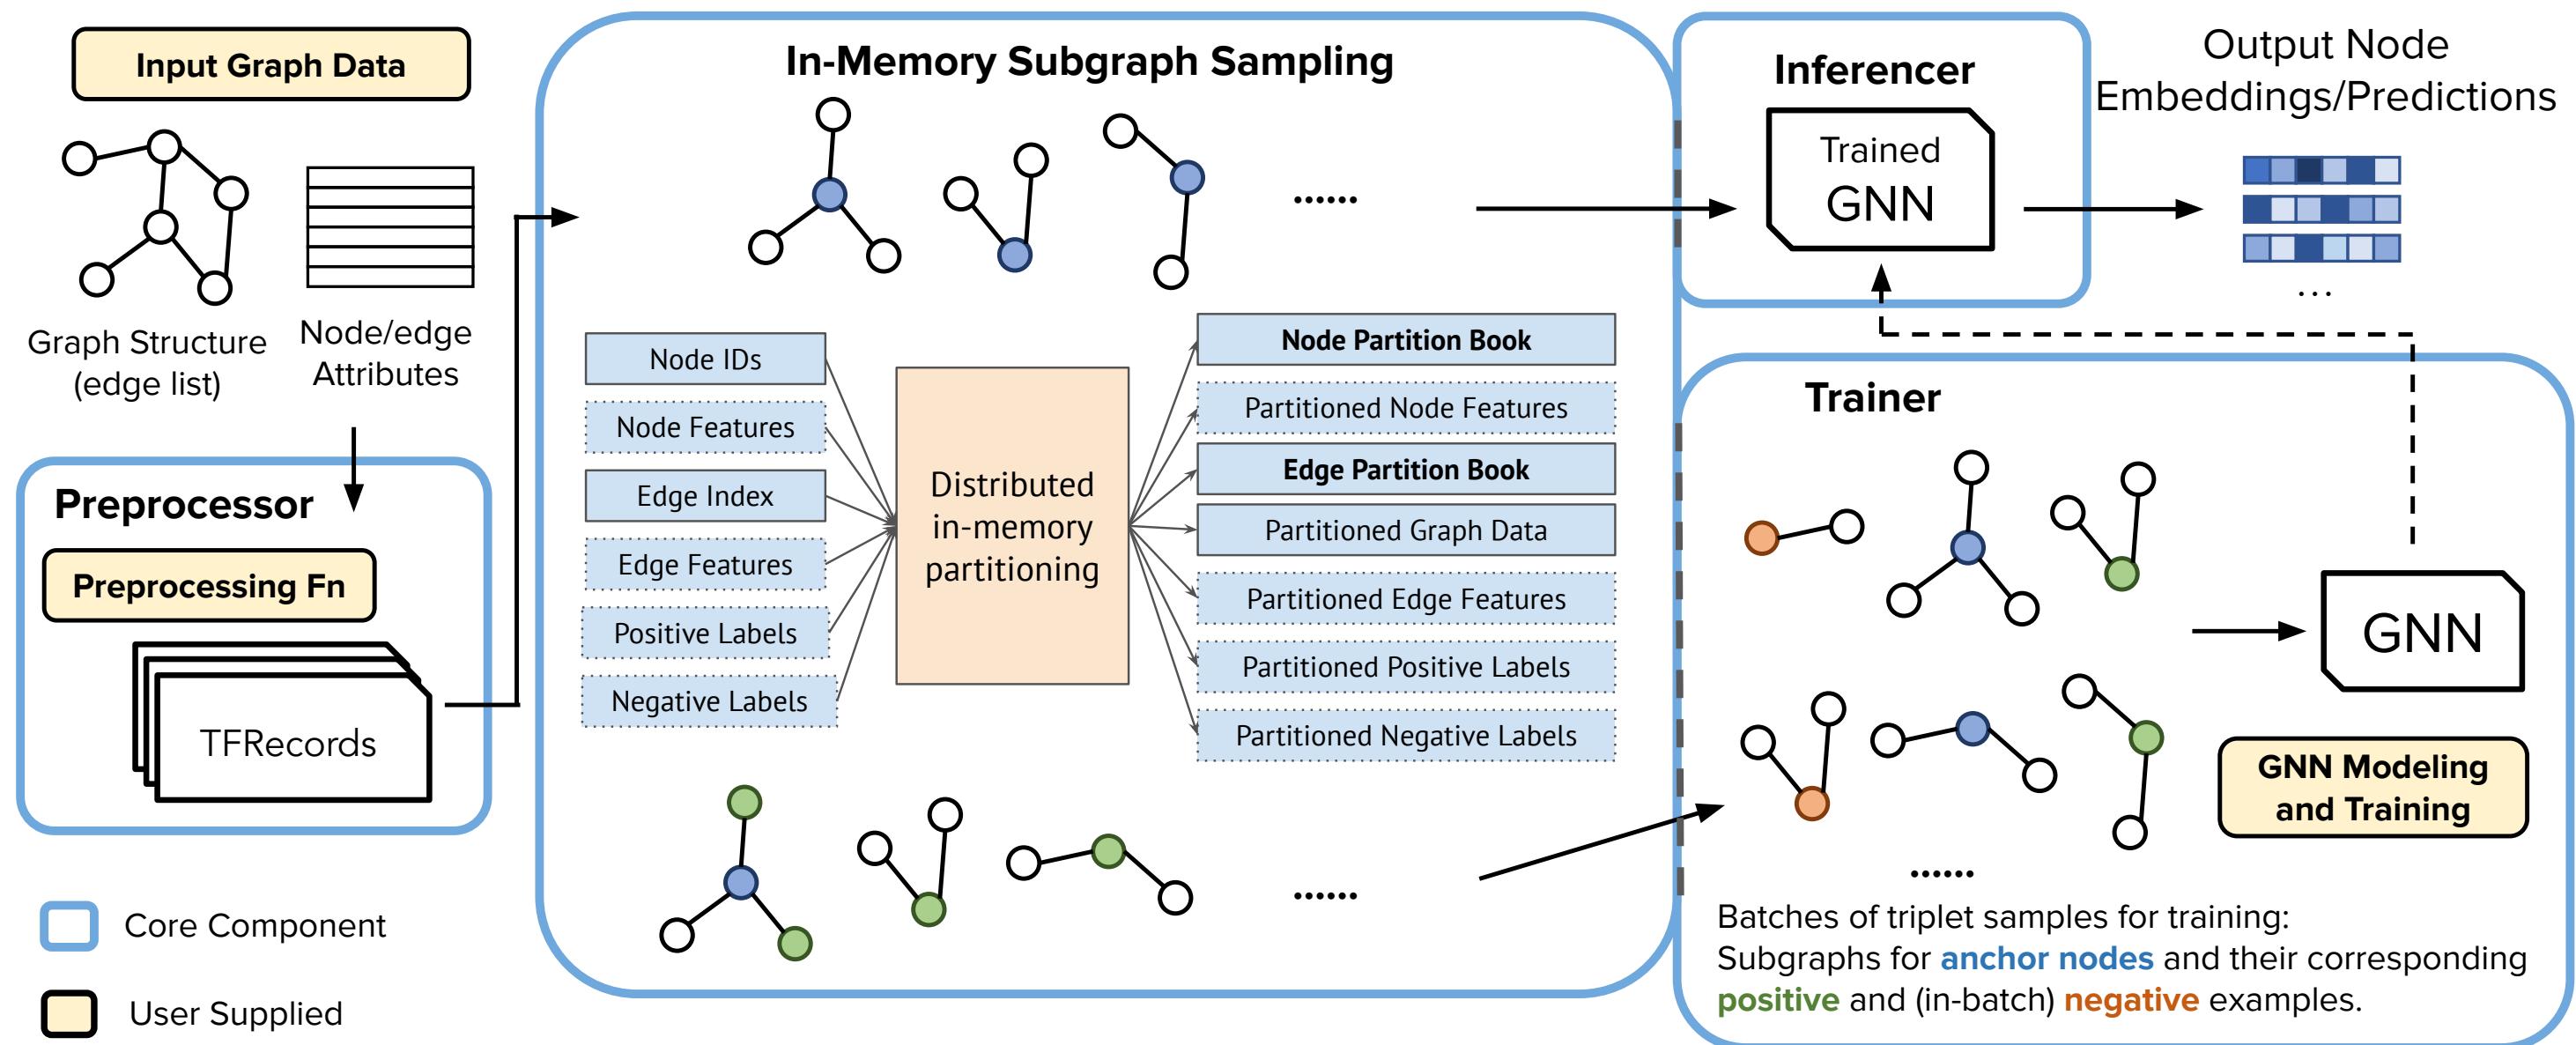
Overview of GiGL

GiGL Components & Orchestration

GiGL components - tabularization



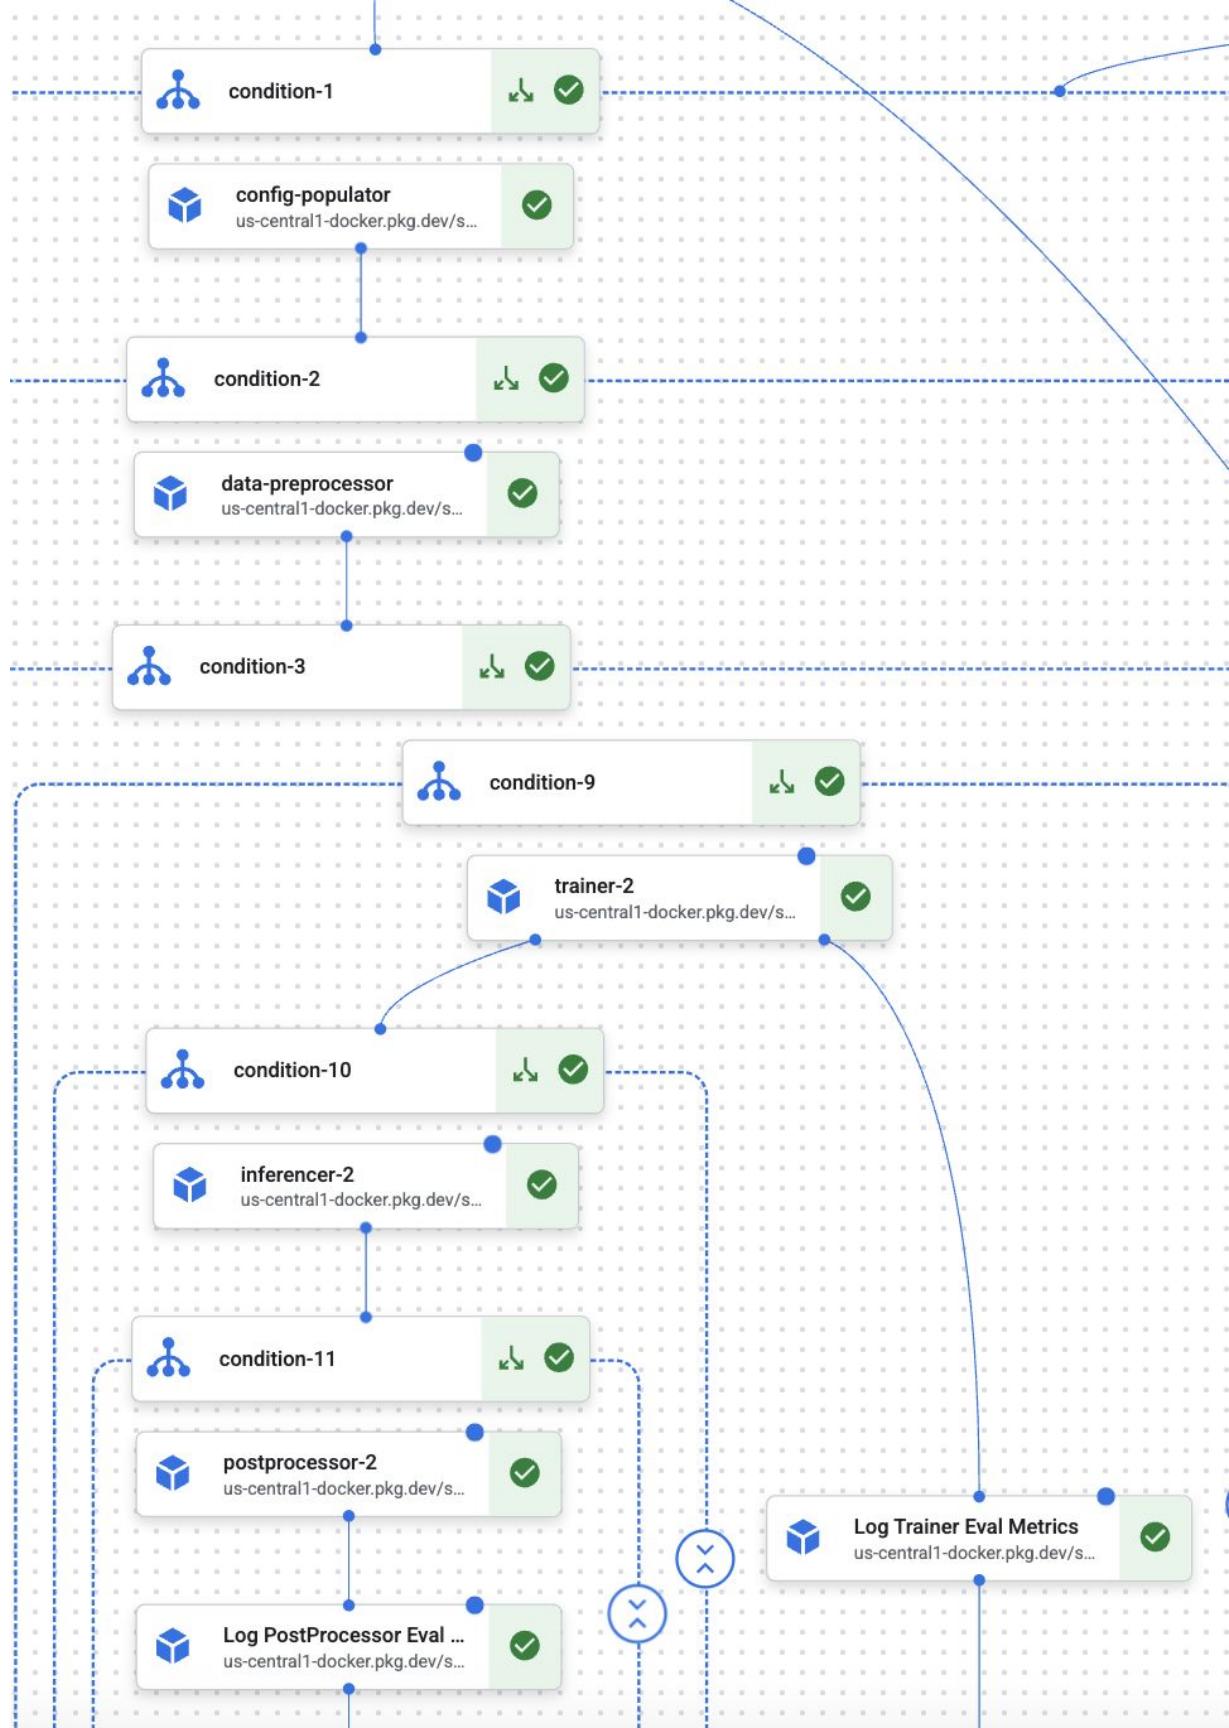
GiGL components - In-memory



GiGL Orchestration

GiGL is orchestrated with Kubeflow pipelines (KFP)

Rightside example: KFP deployed on GCP
Vertex AI pipelines



Where to use GiGL?

- Train graph ML models from small to large scale
- Modeling relational (graph) data with GNNs
 - Target Prediction (node classification, link prediction)
 - Embedding generation (supervised, self-supervised)
 - Yes graphs provides additional information to sequence models
- Recommendation, anomaly detection and more!





Multiple impact verticals in product

Embedding Based Retrieval in Friend Recommendation

Jiahui Shi, Vivek Chaurasiya, Yozhen Liu, Shubham Vij, Yan Wu, Satya Kanduri
 Neil Shah, Peicheng Yu, Nik Srivastava, Lei Shi, Ganesh Venkataraman, Jun Yu
 jshi3,vchaurasiya,yliu2,svij,ywu,satya.kanduri,nshah,peicheng.yu,nsrivastava,lshi3,gvenkataraman,jyu3}@snap.com
 Snap Inc.
 Santa Monica, CA USA

ABSTRACT

Friend recommendation systems in online social and professional networks such as Snapchat helps users find friends and build connections, leading to better user engagement and retention. Traditional friend recommendation systems take advantage of the principle of locality and use graph traversal to retrieve friend candidates, e.g. Friends-of-Friends (FoF). While this approach has been adopted and shown efficacy in companies with large online networks such as LinkedIn and Facebook, it suffers several challenges: (i) discrete graph traversal offers limited reach in cold-start settings, (ii) it is expensive and infeasible in realtime settings beyond 1 or 2 hop requests owing to latency constraints, and (iii) it cannot well-capture the complexity of graph topology or connection strengths, forcing one to resort to other mechanisms to rank and find top- K candidates. In this paper, we proposed a new *Embedding Based Retrieval (EBR)* system for retrieving friend candidates, which complements the traditional FoF retrieval by retrieving candidates beyond 2-hop, and providing a natural way to rank FoF candidates.

1 FRIEND RECOMMENDATION SYSTEMS

In online social and professional networks [2, 6, 19, 20], a user's friends or connections are critical for one's engagement and retention. Research at LinkedIn [36] showed that "members with at least 13 connections from companies other than their current employer are 22.9% faster in transitioning to their next job than those who do not". Friend recommendation can be formulated as a link prediction problem [22], where the goal is to predict the links that are to be formed at timestamp T given a snapshot of a social network at timestamp $T - 1$. However, unlike classical link prediction problems in academic settings, friend recommendation in online networks operate on hundreds of millions or even billions of users, and evaluating link likelihood between every pair of users is computationally infeasible. Thus, in practice, friend recommendation is often formulated as an industrial recommendation problem and follows a typical large-scale recommendation system architecture [4] which consists of two tiers: *Retrieval* and *Ranking*.

Improving Embedding-Based Retrieval in Friend Recommendation with ANN Query Expansion

Pau Perng-Hwa Kung pkung@snapchat.com Snap Palo Alto, CA, United States	Zihao Fan zfan3@snapchat.com Snap Palo Alto, CA, United States	Tong Zhao tzhao@snapchat.com Snap Bellevue, WA, United States
Yozhen Liu yliu2@snapchat.com Snap Santa Monica, CA, United States	Zhixin Lai zlai@snapchat.com Snap Santa Monica, CA, United States	Jiahui Shi jshi3@snapchat.com Snap Palo Alto, CA, United States
Yan Wu ywu@snapchat.com Snap Palo Alto, CA, United States	Jun Yu jyu@snapchat.com Snap Bellevue, WA, United States	Neil Shah nshah@snapchat.com Snap Bellevue, WA, United States
Ganesh Venkataraman gvenkataraman@snapchat.com Snap Palo Alto, CA, United States		

Launches

- GraphSAGE on friend graph (initial launch of EBR)
- Graph definition update: engagement graph
- Graph definition update: graph augmentation
- Model update: GAT
- Loss update: retrieval loss
- Loss update: multi-task training
- Task update: supervised LP with user-defined labels
- EBR scheme update: Stochastic EBR

online improv.

+10%, +11.6%, +7%
+8.9%, +5.8%, +10.8%
+1.0%, +2.1%
+6.5%, +2.92%
+6.2%, +8.8%
+1.2%, +1.3%
+2.2%, +1.8%
10.2%, 13.9%



Multiple impact verticals in product

Launches	online improv.
Spotlight EBR with GNN embeddings	+0.55%, +1.54%, +2.02%
Discover EBR with GNN embeddings	+0.35%, +1.02%, +3.17%

Leveraging GNN embeddings for Content Retrieval

Setup	MRR	HR@1
Base model: 2 layer GAT, sample 15 neighbors per layer	0.39	0.24
+ feature normalization	0.54	0.40
+ further feature eng.	0.59	0.46
+ increase to 3 layers, double training data-	0.64	0.53
+ # of neighbor sampling to 20	0.68	0.57
+ # of neighbor sampling to 30, half hidden dim.	0.70	0.59

Offline iterations for Content Retrieval

Applications	online improv.
Community guideline violation detection	+10.9%, +6.9%, +36%
Bad actor detection	+25.9%, +27.1%
EBR for Display Name Search (DNS)	+6.23%, +4.65%
Social Graph Embeddings in DNS Heavy Ranker	+1.01%, +0.98%, +0.75%
Social Graph Embeddings in Lens Ranking	+0.54%, +0.51%, +0.62%

Multiple applications across Safety, Search and Lens Experience

Setup	# edges	neighbors size	MRR	Precision	Recall
Control (graph embed.)	60B	-	-	0.1523	0.3220
2 layer GAT (v1)	3.36B	15	0.871	0.1138	0.3239
2 layer GAT (v2)	6.46B	15	0.959	0.1505	0.3819
2 layer GAT (v3)	6.46B	40	0.980	0.1523	0.4109

Offline iterations for Ad Ranking

Future of GiGL

- GiGL is actively being developed
- We plan to support and open-source more graph embedding methods and graph tooling in the future
- We welcome open-source contributions!



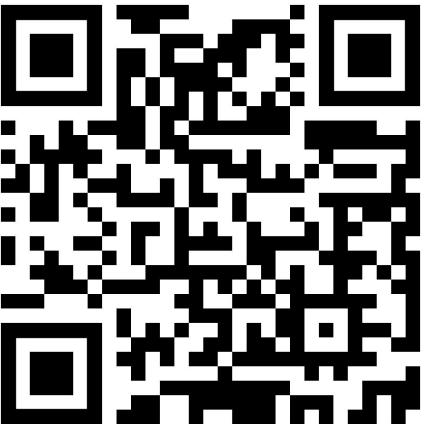


Want to learn more?

Check out our **KDD '25 ADS track paper** online

GiGL is also **open-sourced** for public use;
please enjoy, reach out and contribute

We have ample documentation to get
started building your own pipelines



KDD '25 ADS track paper



Open-source repo (active!)



Documentation and resources

Acknowledgements



Google Labs for Sales Team



GraphLearn-for-PyTorch team



Hands-on with GiGL

Showcase

Let us showcase an End-to-End pipeline run
with a large dataset

([MAG240M](#) - billion scale edges)

Please follow lab instructions

[https://github.com/Snapchat/GiGL/tree/main
/examples/tutorial/KDD_2025](https://github.com/Snapchat/GiGL/tree/main/examples/tutorial/KDD_2025)



Break (30m)

We'll be back at 15:30



Hands-on with GiGL

In-Memory SGS Deep Dive + Customizations

Extending Toy Examples to Large Scale

Resource Config Changes for large-scale runs

Toy Resource Config

```
preprocessor_config:  
  edge_preprocessor_config:  
    num_workers: 1  
    max_num_workers: 2  
    machine_type: "n2-standard-16"  
    disk_size_gb: 300  
  node_preprocessor_config:  
    num_workers: 1  
    max_num_workers: 2  
    machine_type: "n2-standard-16"  
    disk_size_gb: 300  
  trainer_resource_config:  
    vertex_ai_trainer_config:  
      machine_type: "n1-highmem-8"  
      gpu_type: ACCELERATOR_TYPE_UNSPECIFIED # CPU Training  
      gpu_limit: 0 # CPU Training  
      num_replicas: 1  
  inferencer_resource_config:  
    vertex_ai_inferencer_config:  
      machine_type: "n1-highmem-8"  
      gpu_type: ACCELERATOR_TYPE_UNSPECIFIED # CPU Inference  
      gpu_limit: 0 # CPU Inference  
      num_replicas: 1
```

MAG240M Resource Config

```
preprocessor_config:  
  edge_preprocessor_config:  
    num_workers: 1  
    max_num_workers: 256  
    machine_type: "n2d-highmem-64"  
    disk_size_gb: 300  
  node_preprocessor_config:  
    num_workers: 1  
    max_num_workers: 128  
    machine_type: "n2d-highmem-64"  
    disk_size_gb: 300  
  trainer_resource_config:  
    vertex_ai_trainer_config:  
      machine_type: n1-highmem-96  
      gpu_type: NVIDIA_TESLA_T4  
      gpu_limit: 4  
      num_replicas: 20  
  inferencer_resource_config:  
    vertex_ai_inferencer_config:  
      machine_type: n1-highmem-96  
      gpu_type: NVIDIA_TESLA_T4  
      gpu_limit: 4  
      num_replicas: 20
```

Data Preprocessor – Prepare for Pipeline

- We can use the [prepare_for_pipeline](#) function of the [DataPreprocessorConfig](#) class to perform any operation needed prior to running the pipeline. This can include
 - Generating the labeled edge types for supervision edges
 - Filtering existing tables based on some criteria
 - Transforming existing graph tables

```
def prepare_for_pipeline(
    self, applied_task_identifier: AppliedTaskIdentifier
) -> None:
    """
    This function is called at the very start of the pipeline before enumerator and datapreprocessor.
    This function does not return anything. It can be overwritten to perform any operation needed
    before running the pipeline, such as gathering data for node and edge sources

    Args:
        applied_task_identifier (AppliedTaskIdentifier): A unique identifier for the task being run. This is usually
            the job name if orchestrating through GiGL's orchestration logic.
    Returns:
        None
    """
    return None
```

Data Preprocessor – Prepare for Pipeline

- For MAG240M, we use this to generate features for the **author** and **institution** node types by taking the average of their 1-hop neighbors with features.

```
def prepare_for_pipeline(
    self, applied_task_identifier: AppliedTaskIdentifier
) -> None:
    logger.info(
        f"Preparing for pipeline with applied task identifier: {applied_task_identifier}",
    )
    bq_utils = BqUtils(project=self._resource_config.project)

    author_table = (
        f"{self._resource_config.project}.{self._resource_config.temp_assets_bq_dataset_name}."
        + f"{applied_task_identifier}_author_feature_table"
    )

    average_author_query = query_template_compute_average_features.format(
        feature_table=self._paper_table,
        edge_table=self._author_writes_paper_table,
        join_identifier=self._paper_node_type,
        group_by_identifier=self._author_node_type,
        average_feature_query=self._average_feature_query,
    )
    bq_utils.run_query(
        query=average_author_query,
        labels=[],
        destination=author_table,
        write_disposition=WriteDisposition.WRITE_TRUNCATE,
    )
```

```
query_template_compute_average_features = """
WITH joined_table AS (
    SELECT
        *
    FROM `'{feature_table}`` AS feature_table
    JOIN
        `'{edge_table}`` AS edge_table
    ON
        feature_table.{join_identifier} = edge_table.{join_identifier}
)
SELECT
    {group_by_identifier},
    {average_feature_query}
FROM
    joined_table
GROUP BY
    {group_by_identifier};
    """
```

Data Preprocessor – Adding Features

- The `build_ingestion_feature_spec_fn` in our preprocessing config is used to identify the input columns and their types from the source, in this case BigQuery.
- For our toy example, we provide our node identifiers in the `fixed_int_fields` and the features in `fixed_float_fields`, which are the same across both node types.

```
feature_spec_fn = build_ingestion_feature_spec_fn(  
    fixed_int_fields=[node_identifier],  
    fixed_float_fields=self._node_float_feature_list,  
)
```

Data Preprocessor – Adding Features

- For more complex datasets, we can specify any typed combination of fields, and can vary this for each node and edge type
- Each node and edge type can have their own preprocessing functions for their features

```
def build_ingestion_feature_spec_fn(
    fixed_string_fields: Optional[list[str]] = None,
    fixed_string_field_shapes: dict[str, list[int]] = {},
    fixed_float_fields: Optional[list[str]] = None,
    fixed_float_field_shapes: dict[str, list[int]] = {},
    fixed_int_fields: Optional[list[str]] = None,
    fixed_int_field_shapes: dict[str, list[int]] = {},
    varlen_string_fields: Optional[list[str]] = None,
    varlen_float_fields: Optional[list[str]] = None,
    varlen_int_fields: Optional[list[str]] = None,
) -> Callable[[], FeatureSpecDict]:
```

Data Preprocessor – Preprocessing Functions

Passthrough Function

```
def preprocessing_fn(inputs: TFTensorDict) -> TFTensorDict:  
    return inputs
```

Normalizing Function

```
def preprocessing_fn(inputs: TFTensorDict) -> TFTensorDict:  
    outputs = inputs.copy()  
    for feature_name in numerical_features:  
        # outputs[feature] = tft.scale_to_z_score(inputs[feature]) would be used normally,  
        # but this is not executable in a python environment, so we can calculate z-score manually below  
        feat = inputs[feature_name]  
        mean, variance = tf.nn.moments(feat, axes=0)  
        stddev = tf.sqrt(variance)  
        outputs[feature_name] = (feat - mean) / stddev  
    for feature_name in categorical_features:  
        feat = inputs[feature_name]  
        # Extract unique city names  
        unique_count, _ = tf.unique(feat)  
  
        # Create a lookup table mapping each unique city to a unique integer  
        lookup_table = tf.lookup.StaticHashTable(  
            initializer=tf.lookup.KeyValueTensorInitializer(  
                keys=unique_count,  
                values=tf.range(tf.size(unique_count, out_type=tf.int64), dtype=tf.int64)  
            ),  
            default_value=-1 # This is the value returned for any key not in the table  
        )  
  
        # Use the lookup table to convert city names to integers  
        encoded_cities = lookup_table.lookup(feat)  
        outputs[feature_name] = tf.reshape(encoded_cities, tf.shape(feat))  
    return outputs
```

Adding Modeling Levers

Task Definition - Self-Supervised Link Prediction

- Link prediction is specified in the **taskMetadata** field

```
taskMetadata:  
  nodeAnchorBasedLinkPredictionTaskMetadata:  
    supervisionEdgeTypes:  
      - dstNodeType: story  
        relation: to  
        srcNodeType: user
```

- Self-supervised edges can be selected as an argument in **trainerArgs** field

```
trainerConfig:  
  trainerArgs:  
    ssl_positive_label_percentage: "0.7"  
    command: python -m examples.tutorial.KDD_2025.heterogeneous_training
```

Task Definition - Supervised Link Prediction

Alternatively, we can specify positive and negative edges in the [DataPreprocessorConfig](#) class for the link prediction task.

If we have two BigQuery tables which contains positive and negative supervision edges. We can preprocess these edges in `get_edges_preprocessing_spec()`

```
main_edge_ref = BigqueryEdgeDataReference(  
    reference_uri=table,  
    edge_type=edge_type,  
    edge_usage_type=EdgeUsageType.MAIN,  
)
```

Main Edge Reference

```
positive_edge_data_ref = BigqueryEdgeDataReference(  
    reference_uri=self.positive_edge_table,  
    edge_type=labeled_edge_type,  
    edge_usage_type=EdgeUsageType.POSITIVE,  
)  
  
negative_edge_data_ref = BigqueryEdgeDataReference(  
    reference_uri=self.negative_edge_table,  
    edge_type=labeled_edge_type,  
    edge_usage_type=EdgeUsageType.NEGATIVE,  
)
```

Labeled Reference

```
return {  
    main_edge_data_ref: edge_data_preprocessing_spec,  
    positive_edge_data_ref: edge_data_preprocessing_spec,  
    negative_edge_data_ref: edge_data_preprocessing_spec  
}
```

EdgeDataPreprocessingSpec

Task Definition - Supervised Link Prediction

This will populate our YAML at the **preprocessedMetadataURI** with fields specifically for positive and negative edges for that edge type.

These fields are automatically compatible with our training and inference data loaders and will load these labeled edges instead of the self-supervised edges.

```
condensedEdgeTypeToPreprocessedMetadata:  
  '0':  
    mainEdgeInfo:  
      dstNodeIdKey: dst  
      mainEdgeInfo:  
        featureDim: 2  
        featureKeys:  
          - f0  
          - f1  
        schemaUri: ...  
        tfrecordUriPrefix: ...  
    negativeEdgeInfo:  
      featureDim: 0  
      schemaUri: ...  
      tfrecordUriPrefix: ...  
    positiveEdgeInfo:  
      featureDim: 0  
      schemaUri: ...  
      tfrecordUriPrefix: ...  
  ...
```

Task Definition - Node Classification

- We have enabled node classification support for tabularized SGS jobs and are in the process of enabling it for in-memory SGS training and inference.
- Node Classification can be specified in the **taskMetadata** field

```
taskMetadata:  
  nodeBasedTaskMetadata:  
    supervisionNodeTypes:  
      - user
```

Trainer – Hyperparameter Tuning

- Any training loop can be specified by the **command** field of the **trainerConfig**
- Arguments can be provided through **trainerArgs**

```
trainerConfig:  
  trainerArgs:  
    fanout: "[10, 10]"  
    learning_rate: "0.0005"  
    batch_size: "256"  
  command: python -m examples.link_prediction.heterogeneous_training
```

Inferencer – Hyperparameter Tuning

- Similarly, inference loops can be specified by the **command** field of the **inferencerConfig**
- Arguments can be provided through **inferencerArgs**
 - Batch size can alternatively be provided through **inferenceBatchSize**

```
inferencerConfig:  
  inferencerArgs:  
    fanout: "[10, 10]"  
  inferenceBatchSize: 512  
  command: python -m examples.link_prediction.heterogeneous_inference
```

Post Processor

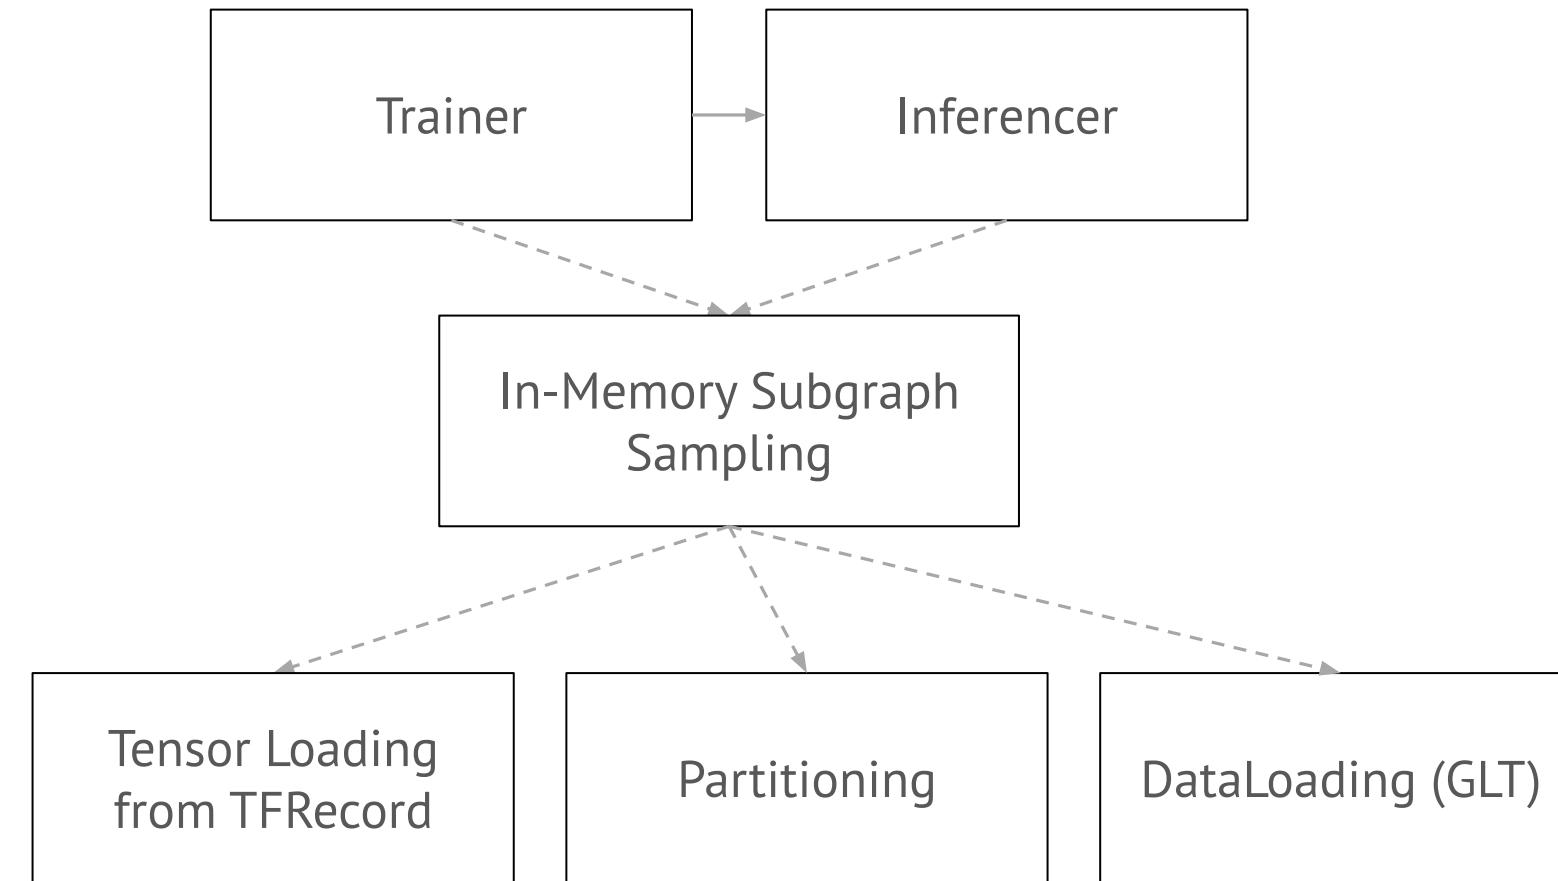
- To allow flexibility for any post-inferencer operations, we also have an **PostProcessor** component which can be optionally specified in the task config
- Handles operations such as asset cleanup, outputting metrics, etc

```
postProcessorConfig:  
  postProcessorArgs: {}  
  postProcessorClsPath: path.to.custom.postProcessor
```

```
class BasePostProcessor(ABC):  
    """  
    Post processor does all operations required after inferencer.  
    Ex. persist inferencer output assets to text files, or run checks on output metrics etc.  
    """  
  
    @abstractmethod  
    def run_post_process(  
        self, gbml_config_pb: gbml_config_pb2.GbmlConfig  
    ) -> Optional[EvalMetricsCollection]:  
        raise NotImplementedError
```

In-Memory SGS Deep Dive and Tuning

In-Memory Subgraph Sampling Overview



In-Memory Subgraph Sampling Overview

```
fanout = [10, 10]
dataset = build_dataset_from_task_config_uri(...)

main_loader = DistABLPLoader(
    dataset=dataset, input_nodes=dataset.train_node_ids, fanout=fanout
)
random_loader = DistNeighborLoader(
    dataset=dataset, input_nodes=dataset.node_ids, fanout=fanout,
)

model: torch.nn.parallel.DistributedDataParallel = ...
loss_fn = ...
optimizer = ...

for batch_idx, (main_data, random_data) in enumerate(zip(main_loader, random_loader)):
    optimizer.zero_grad()
    loss = compute_loss(main_data, random_data, model, loss_fn)
    loss.backward()
    optimizer.step()
```

TFRecord Loading + Partitioning

Dataloading

Training

In-Memory Subgraph Sampling Components

Tensor Loading
from TFRecord

Partitioning

Dataloading

In this step, we load the TFRecords from DataPreprocessor into memory across multiple machines. Each machine will house a subset of the data as torch tensors containing information such as:

- Node Features
- Edge Indices
- Edge Features
- Positive and Hard Negative Labels

In-Memory Subgraph Sampling Components

Tensor Loading
from TFRecord

Partitioning

Dataloading

Two options for how we can load TFRecords:

- **Sequential**: Loads all entities (Node, Edge, Labels) sequentially. Slower, but has smaller peak memory usage. Recommended for cases where memory is a bottleneck
- **Parallel**: Loads all entities (Node, Edge, Labels) in parallel. Faster but larger peak memory usage. Recommended for smaller datasets

Can be determined by setting the **should_load_tensors_in_parallel** flag, defaults to True in our examples.

```
trainerConfig:  
  trainerArgs:  
    | should_load_tensors_in_parallel: True
```

In-Memory Subgraph Sampling Components

Tensor Loading
from TFRecord

Partitioning

Dataloading

In this step, we shuffle the graph data across the machines so that each machine's edges live on the same machine as its source node (if the edge direction is outward) or destination node (if the edge direction is inward)

Generates *partition_books* which store information about what machine each edge index, edge feature, node feature, and labels live on

The partitioned data is subsequently stored in a **DistLinkPredictionDataset** class to be used for data loading.

In-Memory Subgraph Sampling Components

Tensor Loading
from TFRecord

Partitioning

Dataloading

Two options for how we can partition the data

- **Tensor-based Partitioning:** The partition book is a large tensor equal in size to the number of entities (nodes, edges). This has higher memory usage, but results in faster lookup for identifying which rank an item lives on. Recommended for small datasets.

1	0	2	0	3	3	1	2
---	---	---	---	---	---	---	---

Length = Number of Nodes

Partition_book[5] = 3

In-Memory Subgraph Sampling Components

Tensor Loading
from TFRecord

Partitioning

Dataloading

Two options for how we can partition the data

- **Range-based Partitioning:** Keeps track of *ranges* of IDs. This is *much* more memory efficient.
Recommended for larger datasets.

2	4	6	8
---	---	---	---

Length = Number of Machines

Partition_book[5] = 2

In-Memory Subgraph Sampling Components

Tensor Loading
from TFRecord

Partitioning

Dataloading

Two options for how we can partition the data

- **Tensor-based Partitioning:** The partition book is a large tensor equal in size to the number of entities (nodes, edges). This has higher memory usage, but results in faster lookup for identifying which rank an item lives on. Recommended for small datasets.
- **Range-based Partitioning:** Keeps track of *ranges* of IDs. This is *much* more memory efficient. Recommended for larger datasets.

Can be determined by setting the `should_use_range_partitioning`, defaults to True in our examples.

```
trainerConfig:  
  trainerArgs:  
    | should_use_range_partitioning: True
```

In-Memory Subgraph Sampling Components

Tensor Loading
from TFRecord

Partitioning

Dataloading

- Dataloaders produce subgraph samples on the fly using the partitioned data. Sampling occurs asynchronously – subgraph batches are pre-fetched to improve throughput for training/inference.
- The GiGL anchor-based link prediction (ABLP) Dataloader fans out around anchor nodes and labeled nodes when sampling subgraphs.

In-Memory Subgraph Sampling Tuning

- **Sampling_workers_per_process**
 - Number of workers per process for sampling.
 - Can be tuned to speedup sampling operations – defaults to 4.
- **Sampling_worker_shared_channel_size**
 - The shared memory buffer size allocated for the channel during sampling
 - Can be tuned to decrease memory usage from dataloaders – defaults to `4GB`
- **Process_start_gap_seconds**
 - Spaced out data loader initialization across processes to reduce memory spike
 - Default to 0, should be set to 30 in large-scale settings.

```
trainerConfig:  
  trainerArgs:  
    sampling_workers_per_process: "4"  
    sampling_worker_shared_channel_size: "4GB"  
    process_start_gap_seconds: "30"
```

Revisiting our Large-scale run

Let us take a look at our completed
End-to-End pipeline run

Conclusions

You now have:

- A good understanding of GNN scale challenges and how GiGL solves them
- Hands-on experience and familiarity with GiGL and its components
- Ability to train your own GNNs at large-scale
- We're excited to see how GiGL can be leveraged in your use case!



FAQ

What scale does GiGL support? Internally we have tested GiGL up to $1e9$ nodes and $1e10$ edges.

What's the cost of GiGL? Our internal inference pipelines cost $<\$100$ for our scale. This however will depend on your GCP cost, and feature size.

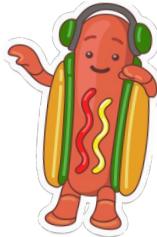
Can GiGL handle temporal graphs? No, we find that Graph Snapshots work well, sequence features can be embedded as node/edge features.

Can GiGL handle real time inference? GiGL can be customized to handle real-time inference.

Does GiGL support graph methods other GNNs? We will be introducing more types of Graph ML methods in the future (ex. Graph embedding based methods etc.)

Can you integrate GiGL with graph databases? We find that in-memory solutions work best for scale, and we're internally moving towards that.

Thank you for attending!



For more questions related to GiGL, please
check the resources below.

[GiGL documentation](#)



[Open source Github repo](#)

